




Kubernetes

Part 2



Kubernetes





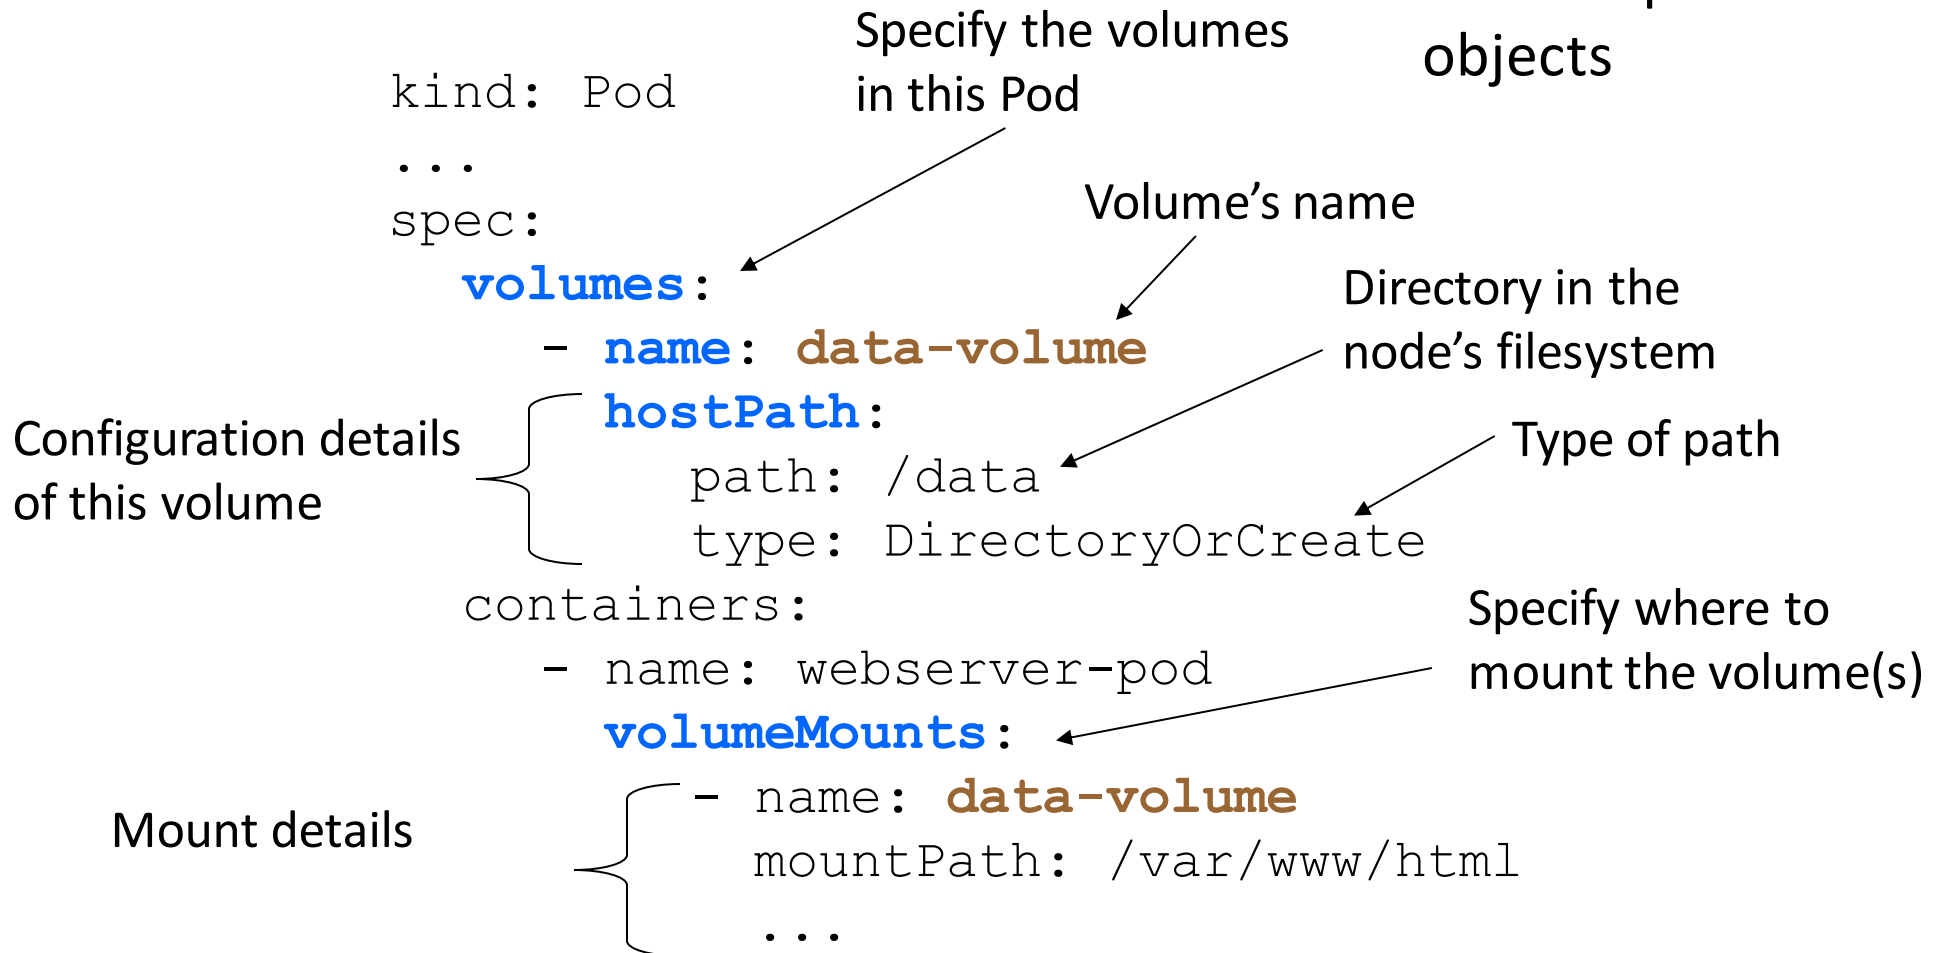
Volumes

- Volumes are storage that are shared by containers in a Pod
 - Allocated by the Pod, usually a shared directory in the Pod
 - Not visible outside of Pod
- Tied to the lifecycle of a Pod viz. its removed when the Pod is delete
 - Unlike Docker volumes where they are durable
- Different types of volumes
 - Eg. hostPath, NFS, iSCSI, fibre channel, empty directory, etc.
- **hostPath** and **emptyDir** type is good for sharing data between containers in a Pod
 - Eg. The example of file puller and web server



Defining a Volume

Same syntax for creating for Pod templates in deployment objects





Using ConfigMaps


Injecting as environment variables

```
containers:
  ...
  env:
    - name: DB_NAME
      valueFrom:
        configMapKeyRef:
          name: myapp-config
          key: db_name
    - name: DB_HOST:
      valueFrom:
        configMapKeyRef:
          name: myapp-config
          key: db_host
```

Mounting as a volume


```
volumes:
  - name: config-volume
    configMap:
      name: myapp-config

containers:
  ...
  volumeMounts:
    - name: config-config
      mountPath: /config
```





Kubernetes





Persistent Storage

- Kubernetes can dynamically provision storage
 - Eg. User ask for 50GB volume to caching images
- Kubernetes allows storage to be either statically or dynamically provisioned
 - Static provision - an administrator will need to first provision the storage manually
 - Dynamic provision - the user describes the type of storage that is required; Kubernetes will attempt to provision based on the user's requirements
- Once a persistent storage has been allocated and claimed/reserved, a Pod can mount the volume like any regular volume
- Persistent volumes lifecycle are not tied to the Pod's lifecycle
 - Unlike volumes, persistent volumes will not be deleted when a Pod is deleted
 - This behaviour can be configured

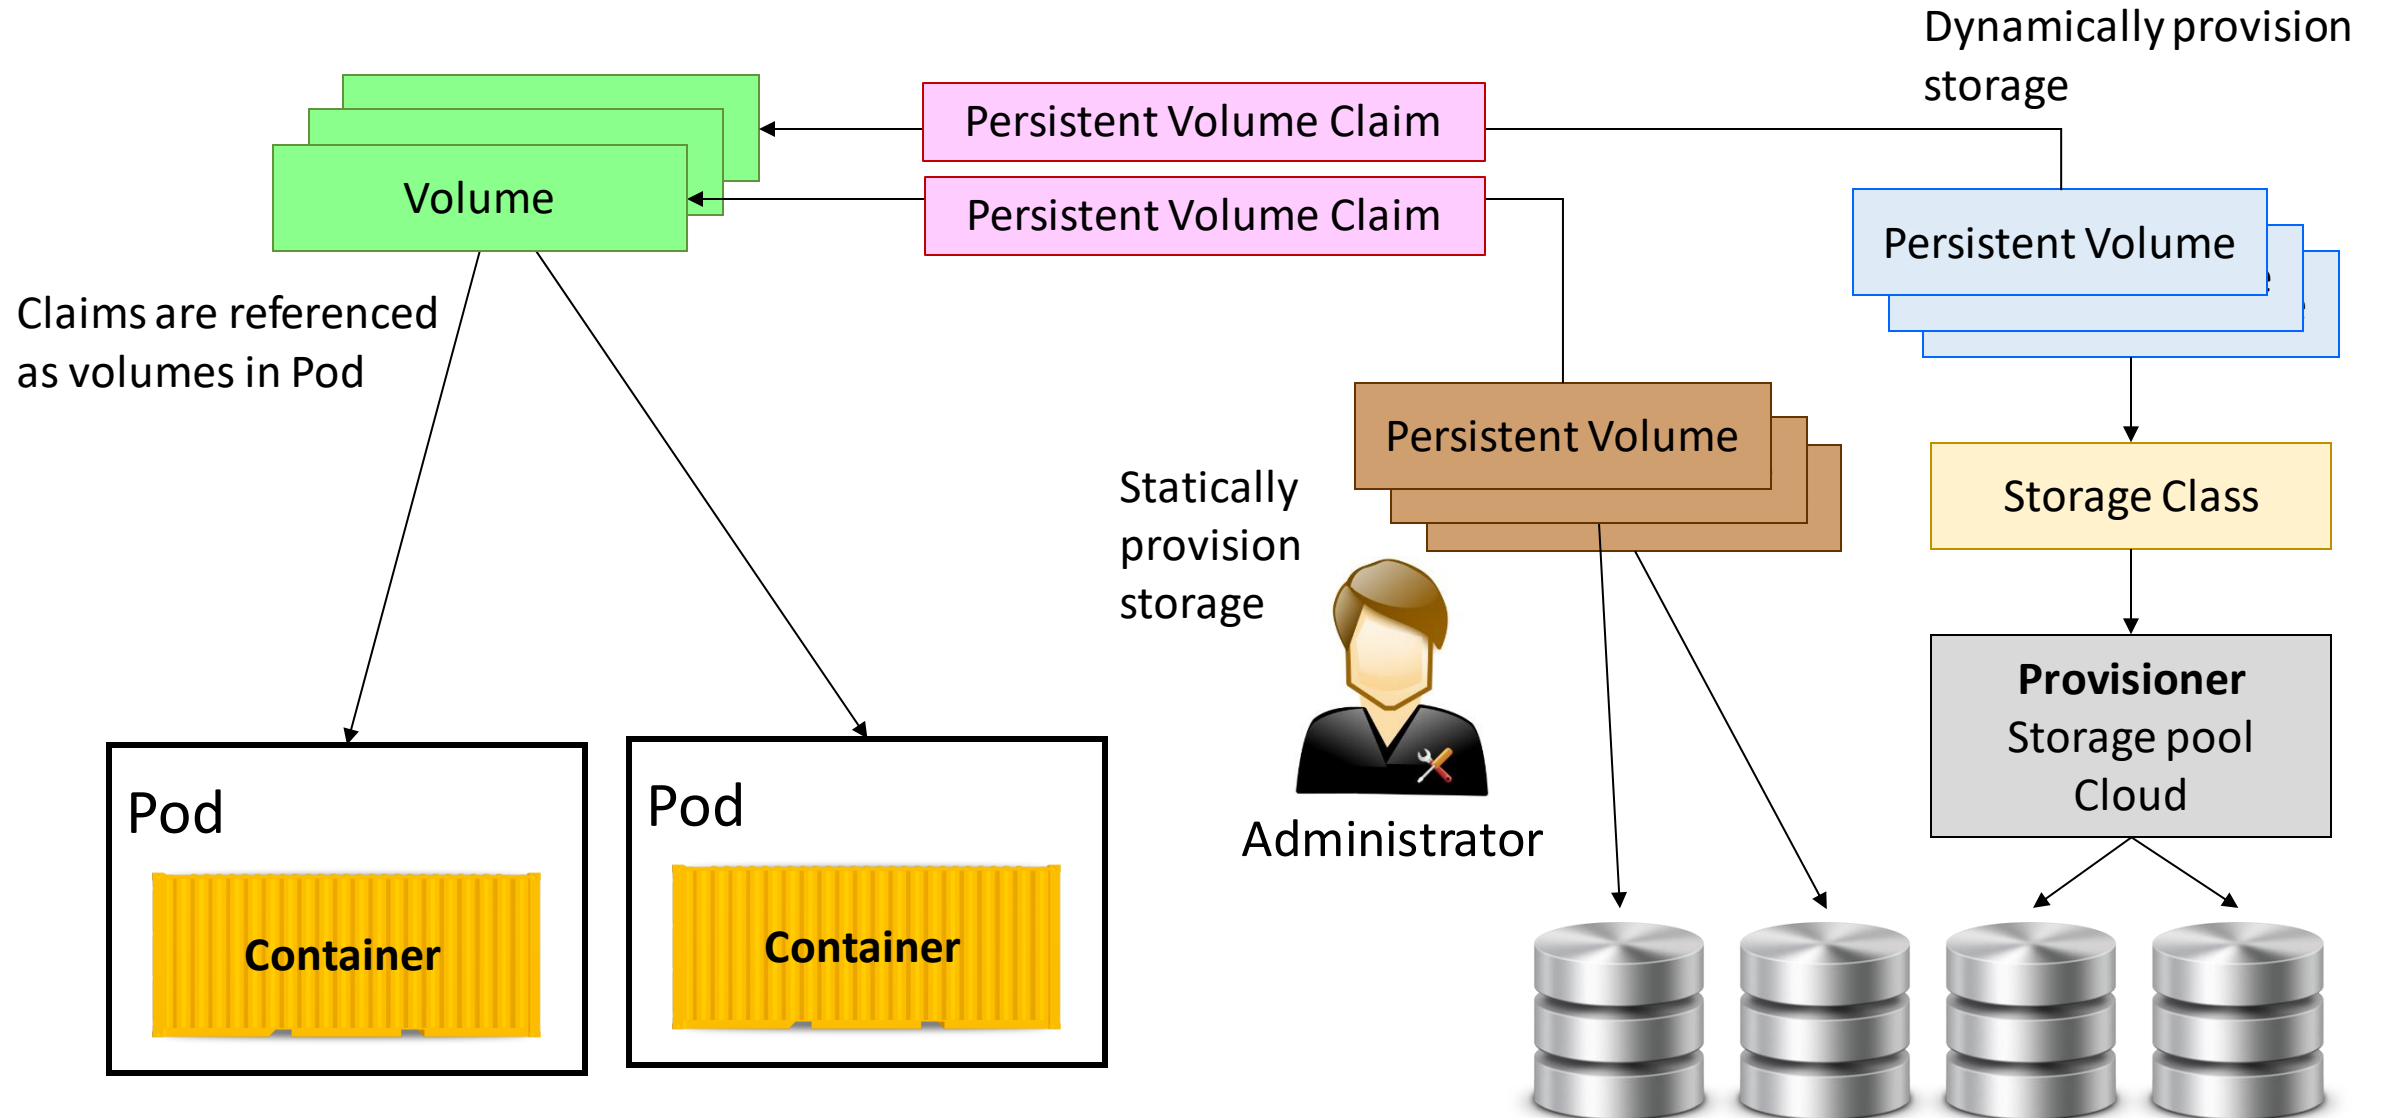


Key Concepts

- Storage class - a type of storage
 - Who the provisioner, storage specific details, retention policy, etc.
- Persistent volume - the actual storage
 - A piece of storage provisioned by an administrator or thru storage class
 - Supports may different storage type
 - AWS EBS, Azure File Service, Cinder, fibre channel, GCP Disk, NFS, etc.
 - Different type of access mode - exclusive or shared
- Persistent Volume claim - when a persistent volume has been allocated for use, the volume is staid to be claimed



Persistent Volume





Static vs Dynamic

Static

- Administrator has to manually allocate storage and map it to a persistent volume
- Users can then claim this volume

Dynamic

- When Kubernetes tries to resolve a claim and the persistent volume is unavailable
- It looks for a storage class that best matches the request storage
- Dynamically creates the persistent volume using the provisioner



Defining a Persistent Volume Claim

apiVersion: v1

kind: PersistentVolumeClaim

meta-data:

name: myapp-pvc

annotations:

volume.beta.kubernetes.io/storage-provisioner: "provisioner"

spec:

accessModes:

- ReadWriteOnly

resources:

requests:

storage: 5Gi

storageClassName: standard

Specify the provisioner that will provisions the storage class

kubectl get storageclass
for a list of provisioners



Mounting a Persistent Volume


```
apiVersion: v1
kind: Pod
meta-data:
  name: myapp
```

```
spec:
```

```
  volumes:
```

```
    - name: data-volume
      persistentVolumeClaim:
        claimName: myapp-pvc
```

Specify the claim name



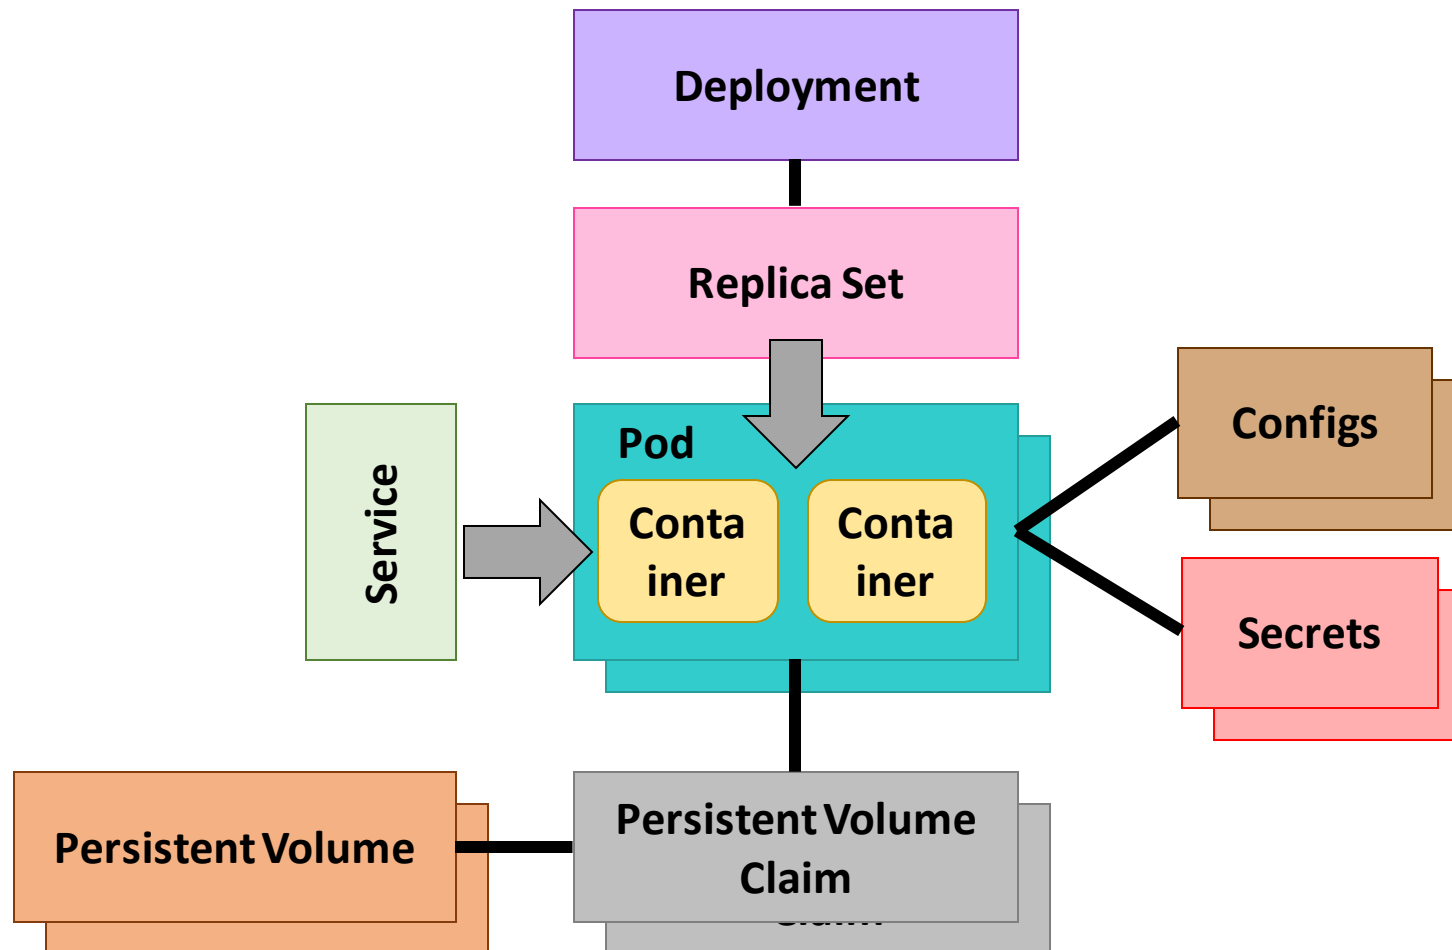
```
  containers:
```

```
    - name: myapp
      volumeMounts:
        - mountPath: /app/public
          name: myapp-pvc
```

```
    ...
```



Kubernetes





Persistence Volume Management

- Display persistence volume detail
 - Persistence volume - `kubectl get pv`
 - Persistence volume claim - `kubectl get pvc`
 - Storage classes - `kubectl get sc`
- Delete persistence volume

```
kubectl delete pvc <name>
```

```
kubectl delete pv <name>
```



Load Balancer and Ingress

- By default services are allocated a cluster IP
 - Only accessible within the cluster
- Load balancer exposes the service to the public
 - Accessible from outside of the cluster
 - Load balancer will redirect the request to pods based on its routing policy
 - Another way to allow external access is via node port
- Load balancer are resources that are provisioned from the underlying cloud platform
 - May have more features that you require
 - Also cost more



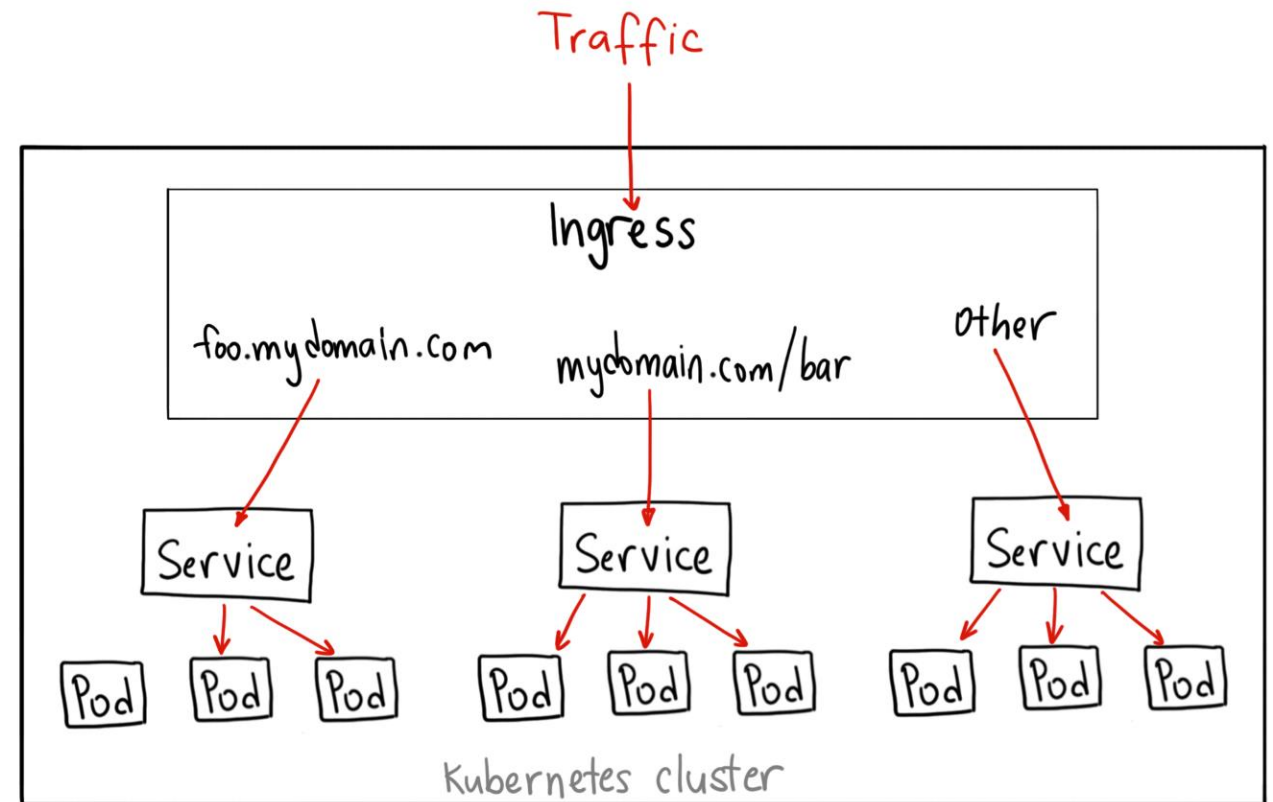
Load Balancer and Ingress

- An ingress is a load balancer but its is provisioned as a service (with pods) inside a Kubernetes cluster
 - It's a resource running in Kubernetes
- Typically less feature and potentially cheaper
 - However you will need to manage it
- NGINX Ingress controller is a popular ingress controller
 - Deploys NGINX as Ingress
 - <https://github.com/kubernetes/ingress-nginx>



Ingress

- Application layer (L7) router that sits in front of multiple services
- Define a set of routing rules on how services are access externally
 - Eg. 2 services, one for search one for checkout. Might map to /search and /checkout
- Rules are applied to ingress controllers which performs the actual routing
 - Controllers might be a cloud provider's load balancer or Nginx reverse-proxy





Defining an Ingress

apiVersion: extensions/v1beta1

kind: Ingress

metadata:

name: myapp

annotations:

Used to configure NGINX
ingress controller

nginx.ingress.kubernetes.io/rewrite-target: "/"
nginx.ingress.kubernetes.io/ssl-redirect: "false"

Change/rewrite a matched resource to
its root e.g /hello to /

spec:

backend:

Default backend if no
rule matches

serviceName: landing
servicePort: 8080

rules:

- http:

paths:

One or more of these rules to
specify which services to
handle what resource

- path: /hello
backend:

serviceName: myapp
servicePort: 8080

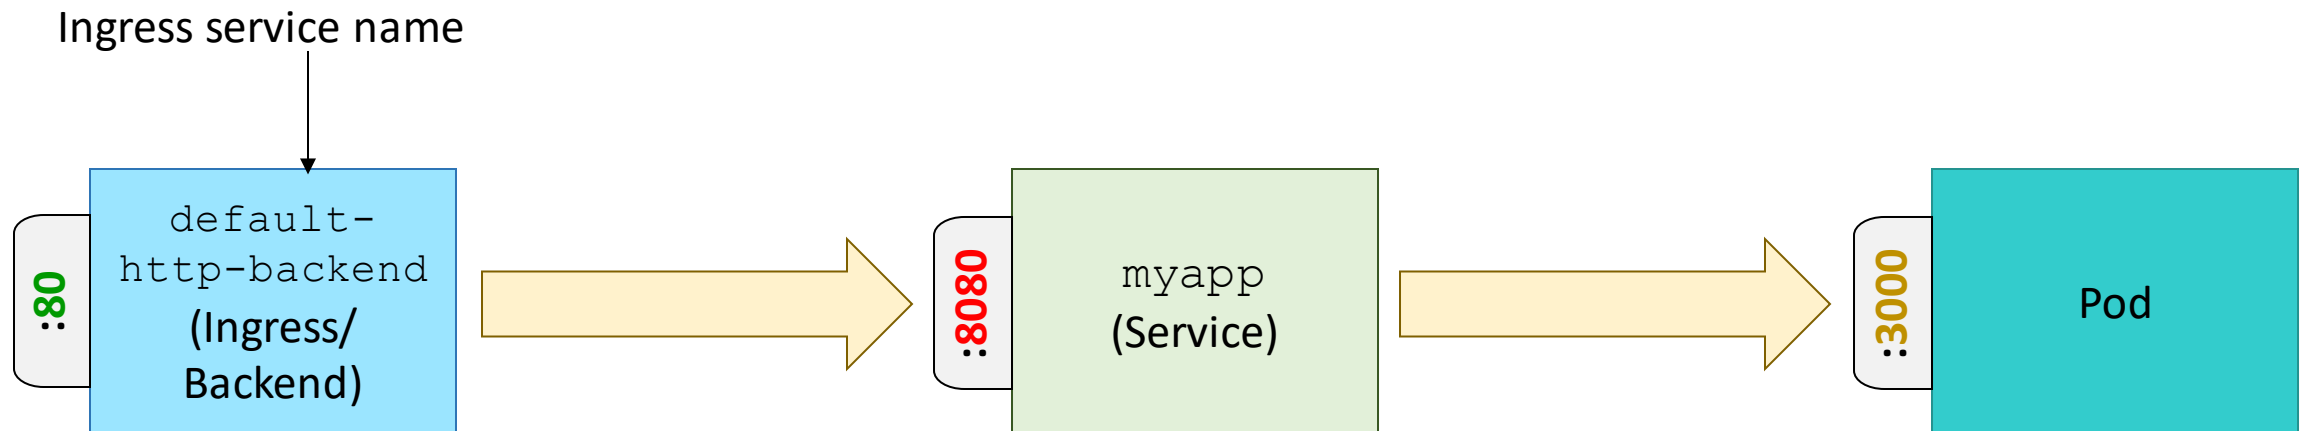


Ingress Ports

```
kind: Ingress
spec:
  rules:
    http:
      paths:
        - backend:
            serviceName: mysvc
            servicePort: 8088
```

```
kind: Service
spec:
  ports:
    - port: 8080
      targetPort: 3000
```

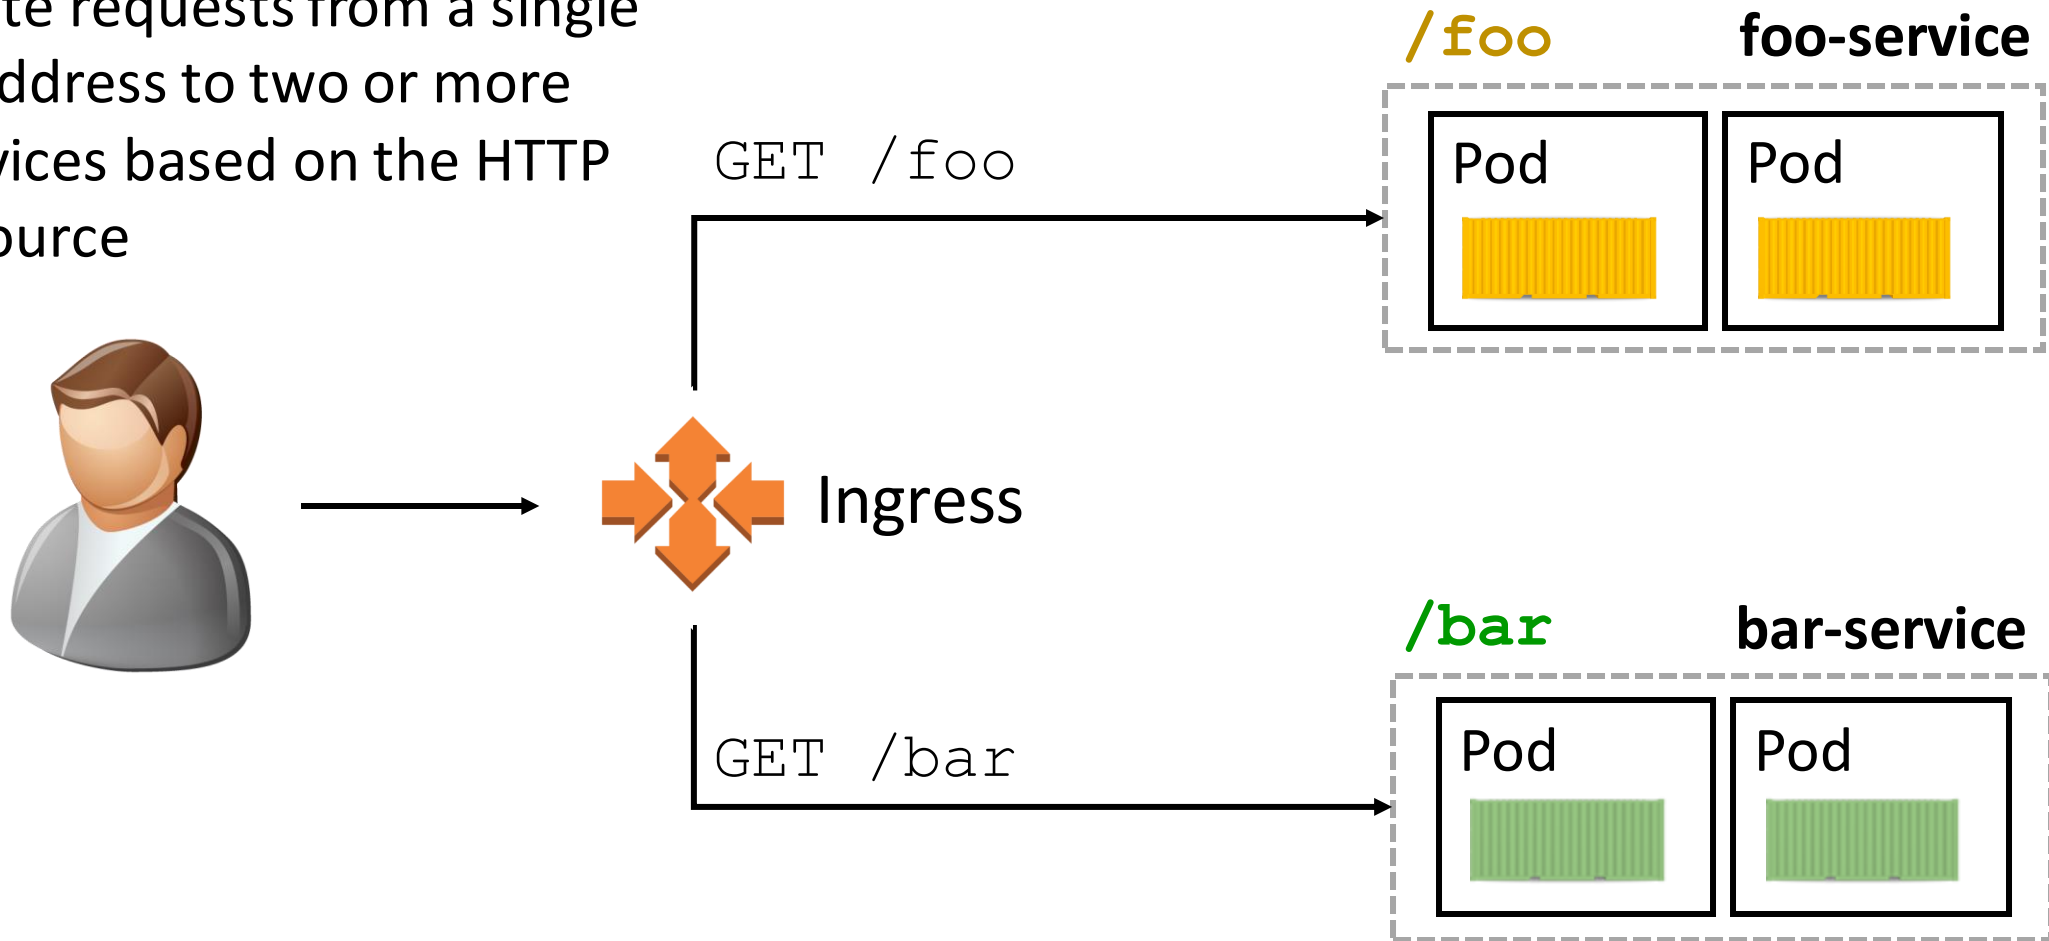
```
kind: Deployment
spec:
  containers:
    ports:
      - containerPort: 3000
```





Ingress - Fan Out

Route requests from a single IP address to two or more services based on the HTTP resource





Ingress Fan Out Example

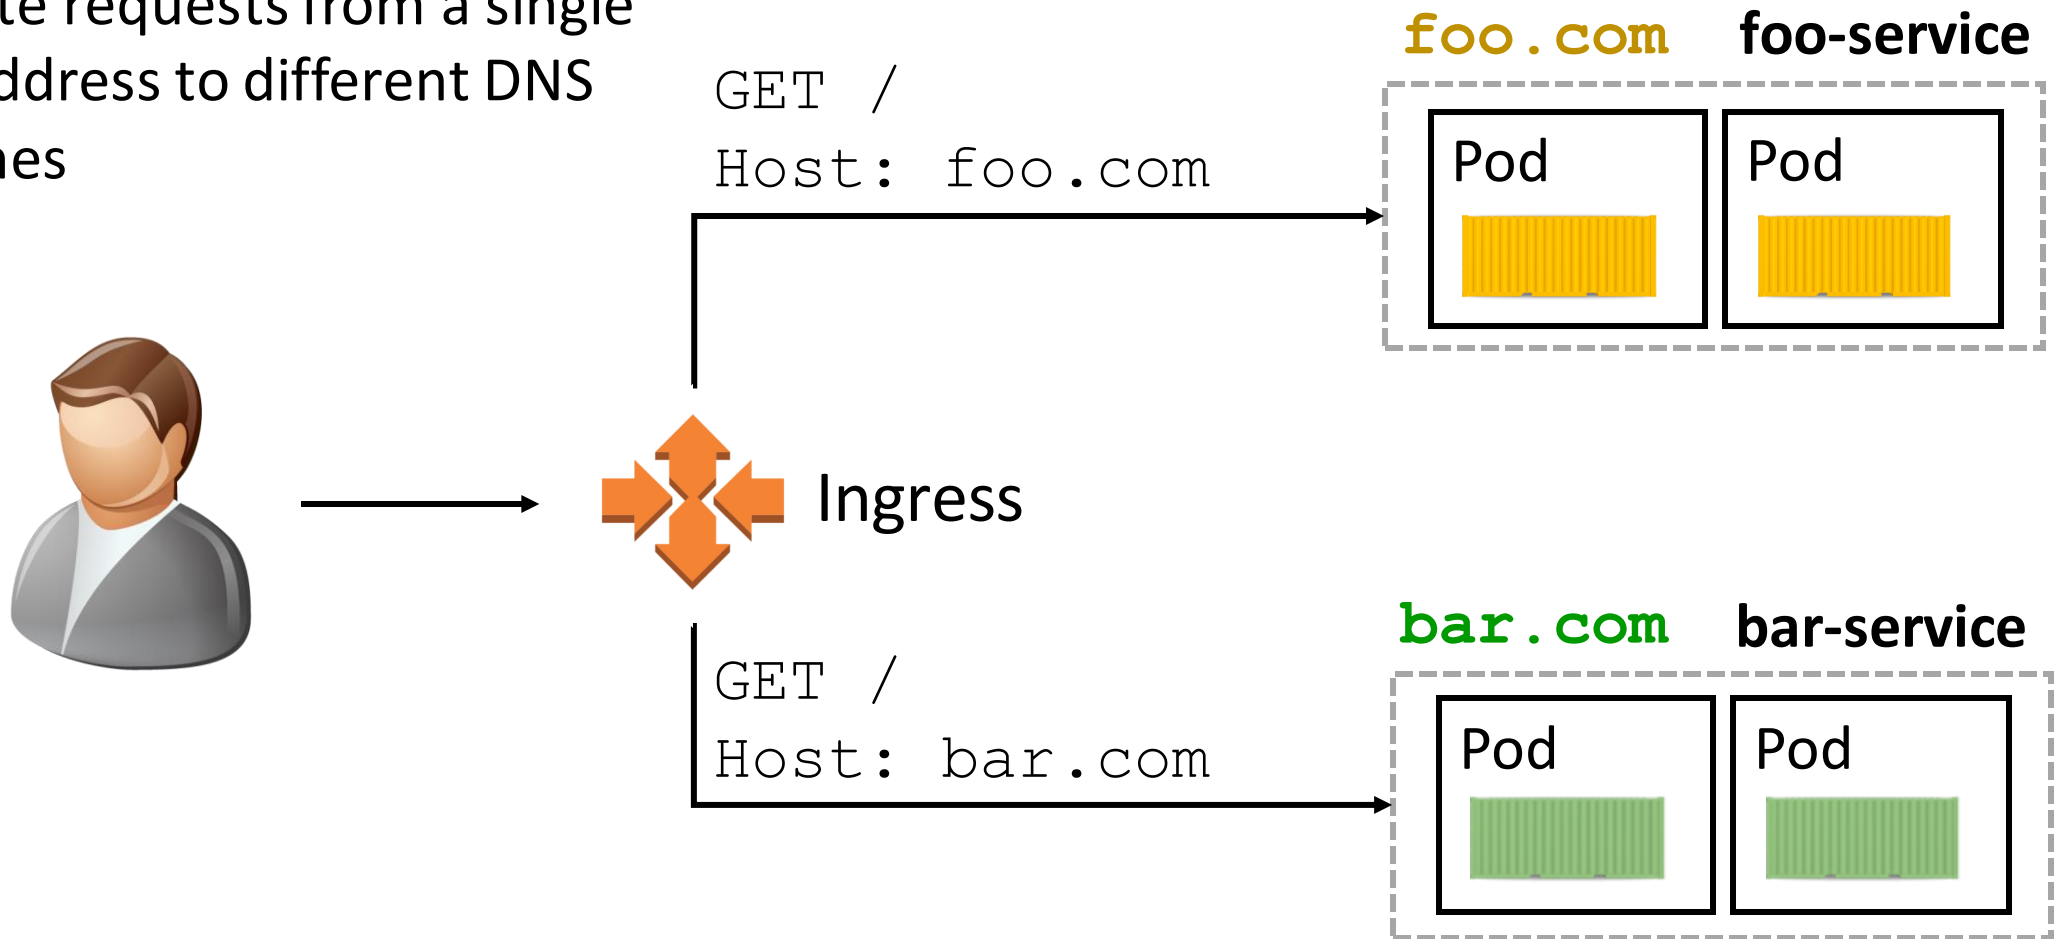
```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  rules:
    - http:
        paths:
          - path: /foo
            backend:
              serviceName: foo-service
              servicePort: 8000
          - path: /bar
            backend:
              serviceName: bar-service
              servicePort: 8001
```

Request is routed to these
2 services depending on
the URI
Default service not shown



Ingress - Virtual Host

Route requests from a single IP address to different DNS names





Ingress Virtual Host Example

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
```

```
  rules:
    - host: foo.com
      http:
        paths:
          - backend:
              serviceName: foo-service
              servicePort: 80
    - host: bar.com
      http:
        paths:
          - backend:
              serviceName: bar-service
              servicePort: 80
```

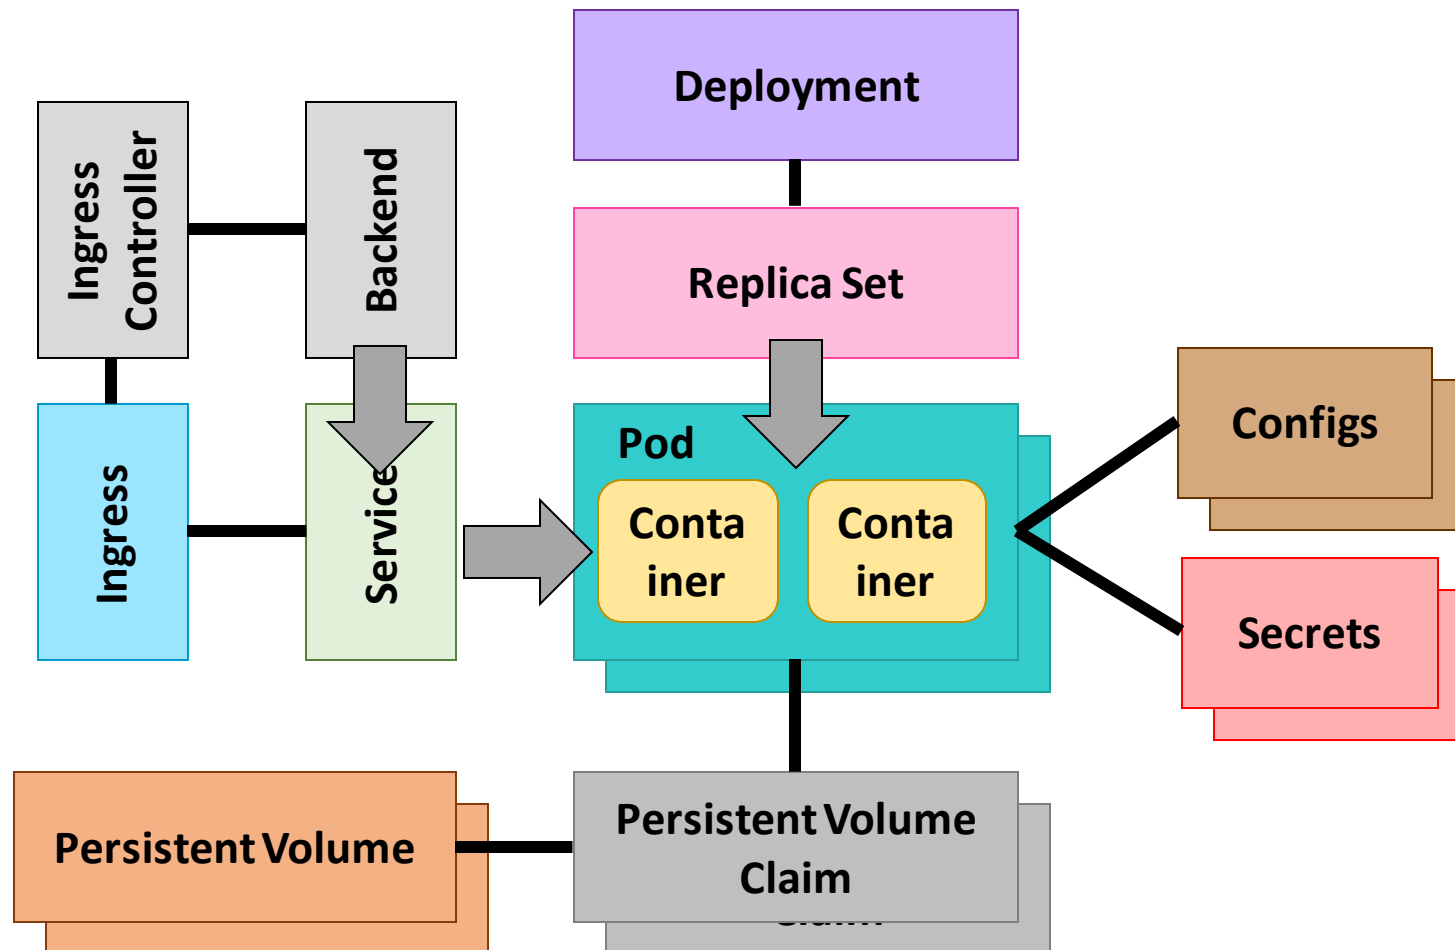
Diagram illustrating the Ingress Virtual Host configuration. The configuration is divided into two sections, each representing a different host:

- foo.com host:** This section defines a rule for the host `foo.com`. It specifies an `http` path that routes requests to the `foo-service` on `port 80`.
- bar.com host:** This section defines a rule for the host `bar.com`. It specifies an `http` path that routes requests to the `bar-service` on `port 80`.

Request is routed to these 2 services depending on the Host attribute. Since the request is to host without specifying a port, the service must expose port 80




Kubernetes





Scaling



- Scaling is the capability of the system to handle more workload by provisioning more resources
- Two types of scaling
 - Horizontal scaling - scales by provision more Pods
 - Applications must be stateless allowing the ingress controller to route the request to any Pod
 - Vertical scaling - scaling by giving the application more resources
 - Application must be able to utilize the extra resources eg. more vCPUs or memory



Why Scale?

- Efficient use of resources
 - Ensure that the actual usage is on parity with the current usage
- Dynamically respond to workload fluctuation
 - Elasticity - providing an agreed on SLA
- Cost optimization
 - Pay only what you use



Horizontal Manual Scaling

- Types of scaling
 - Manual
 - Automatic - Horizontal Pod Autoscaler
- Use `kubectl` to scale up or down

```
kubectl scale --replicas <number> deployment <deployment>
```



Horizontal Pod Autoscaler

- HPA scales a deployment based on one or more metrics
 - Eg. trigger scaling when CPU utilization breaches 80%
 - Metrics to scale the Pods can be
 - Build in metrics , custom metrics, external metrics
- HPA runs a control loop - runs every 30 seconds (default)
 - Queries metrics server
 - Match that against the specified threshold
 - Updates the number of replicas in a deployment if required to meet the load
 - Deployment would then perform the scaling (in or out)
- Reduces cluster size if utilization is low for a period of time
 - Scaling in

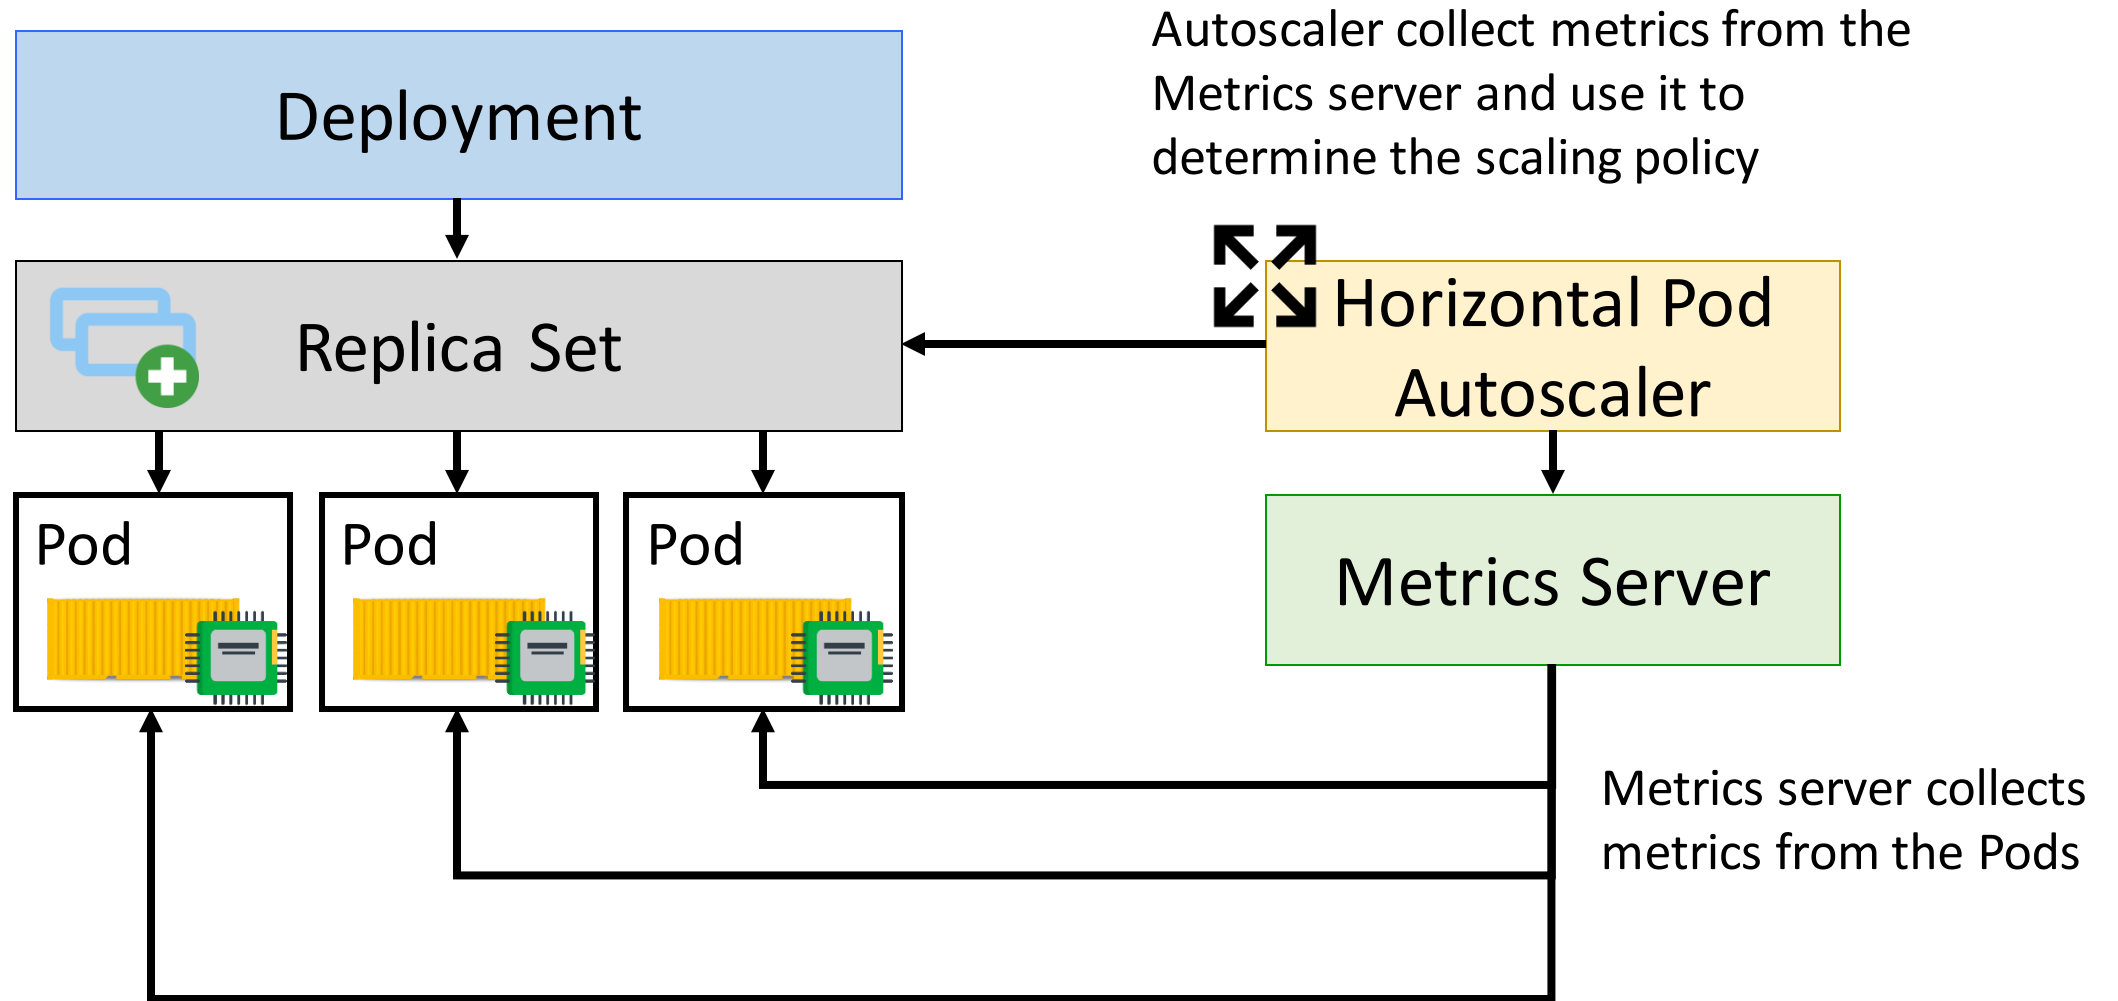


Setting Pod Request

- Horizontal scaler scales a Pod by determining if the Pod has breached a certain threshold
 - For memory and CPU
- Set the request for CPU and memory
 - Specify the minimum amount of compute resources required
- Resource type
 - CPU - measured in CPU units eg 100m is 100 millicores
 - 1 CPU in Kubernetes == 1 vCPU, Core, vCore, Hyperthread
 - Memory - 16M



Horizontal Pod Autoscaler





Requesting Resources

```
apiVersion: v1
kind: Pod
metadata:
  name: myapp
spec:
  containers:
  - name: myapp
    image: myapp:sha256:...
```

HPA only looks
at the CPU

```
resources:
```

```
requests:
```

```
  cpu: 100m
```

```
  memory: 16M
```

```
limits:
```

```
  memory: 32M
```

```
...
```

} Request the minimum amount of
compute resources

→ Describe the maximum amount
of compute resources required



Defining a Horizontal Pod Autoscaler

```
apiVersion: autoscaling/v1
kind: HorizontalPodAutoScaler
metadata:
```

```
  name: myapp
```

```
spec:
```

```
  minReplicas: 3
```

```
  maxReplicas: 8
```

```
  targetCPUUtilizationPercentage: 80
```

```
  scaleTargetRef:
```

```
    apiVersion: apps/v1
```

```
    kind: Deployment
```

```
    name: myapp
```

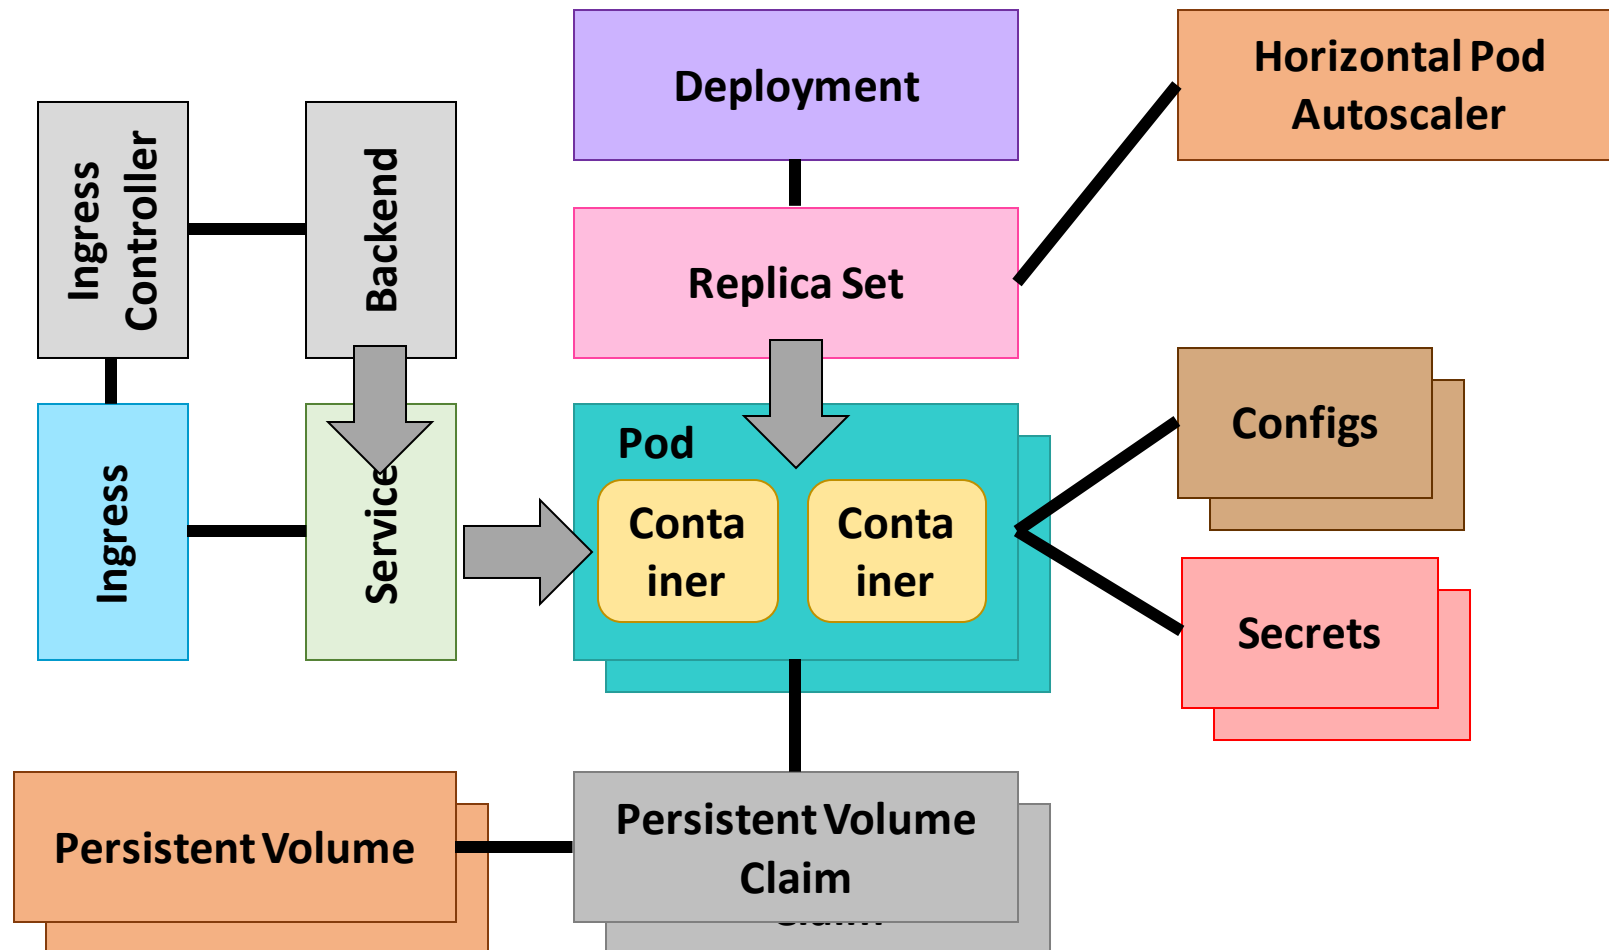
Minimum and maximum number of replicas. Since the HPA is managing the replica set, this setting takes precedence over the deployment setting

The deployment that this HPA is targeting

Percentage of the CPU utilization over all the Pods



Kubernetes





Appendix



Managing Context

- For grouping access parameters under a common name
 - Like a profile
 - Set the namespace, do not need the `-n` option

- Create a context

```
kubectl config set-context <context_name> --namespace=<name> \
    --cluster=<cluster_name> --user=<user_name>
```

- View current contexts

```
kubectl config view
```

- Use a context

```
kubectl config use-context <context_name>
```