



# Istio

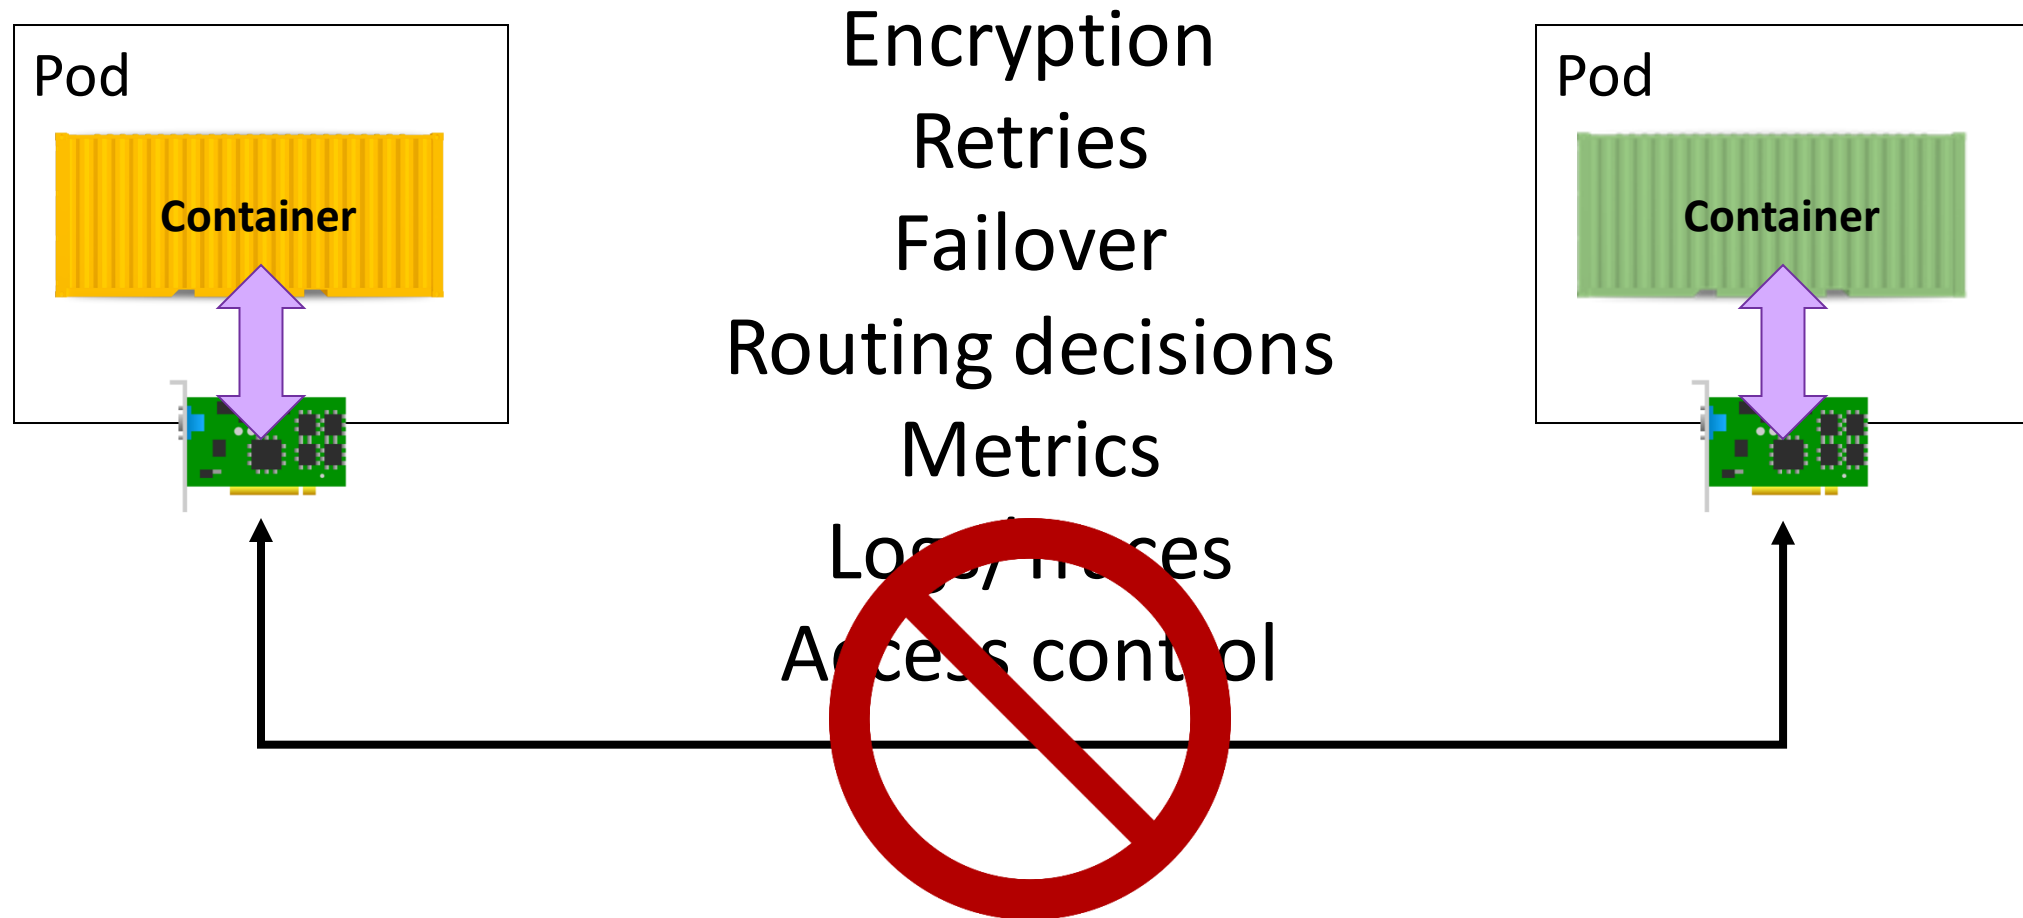


# Microservices and Service Mesh

- An architectural style that decompose an application into loosely coupled services
- Each service implements a single functional concern
- Services communicate with each other over the network
- A collection of these services fulfil a request
  - Service Mesh
- Services need to discover other services
- Service authentication and access control
- Dealing with network issues
- Collecting logs and metrics
- Ability to trace a request as it is fulfilled by multiple services



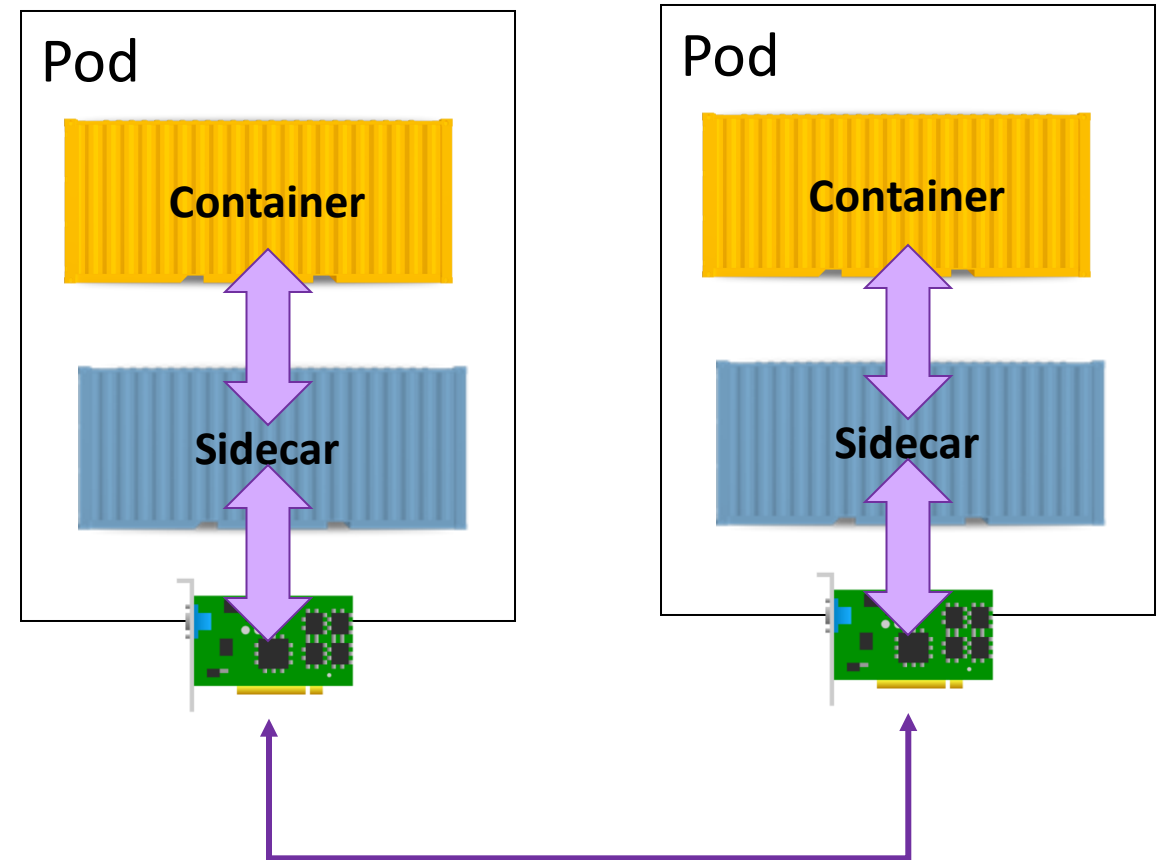
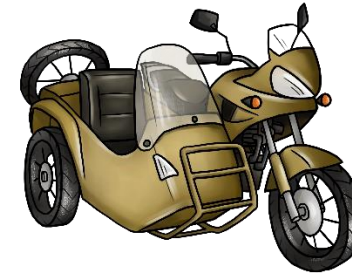
# Microservices Communications





# Sidecar

- Bundle observability, tracing, load balancing, etc. into a separate container
- Inject the container (sidecar) into every Pod
  - Sidecars are part of the infrastructure
- All communications in/out of the Pod is proxied across the sidecar
- Communicate with sidecar for configuration and data collection



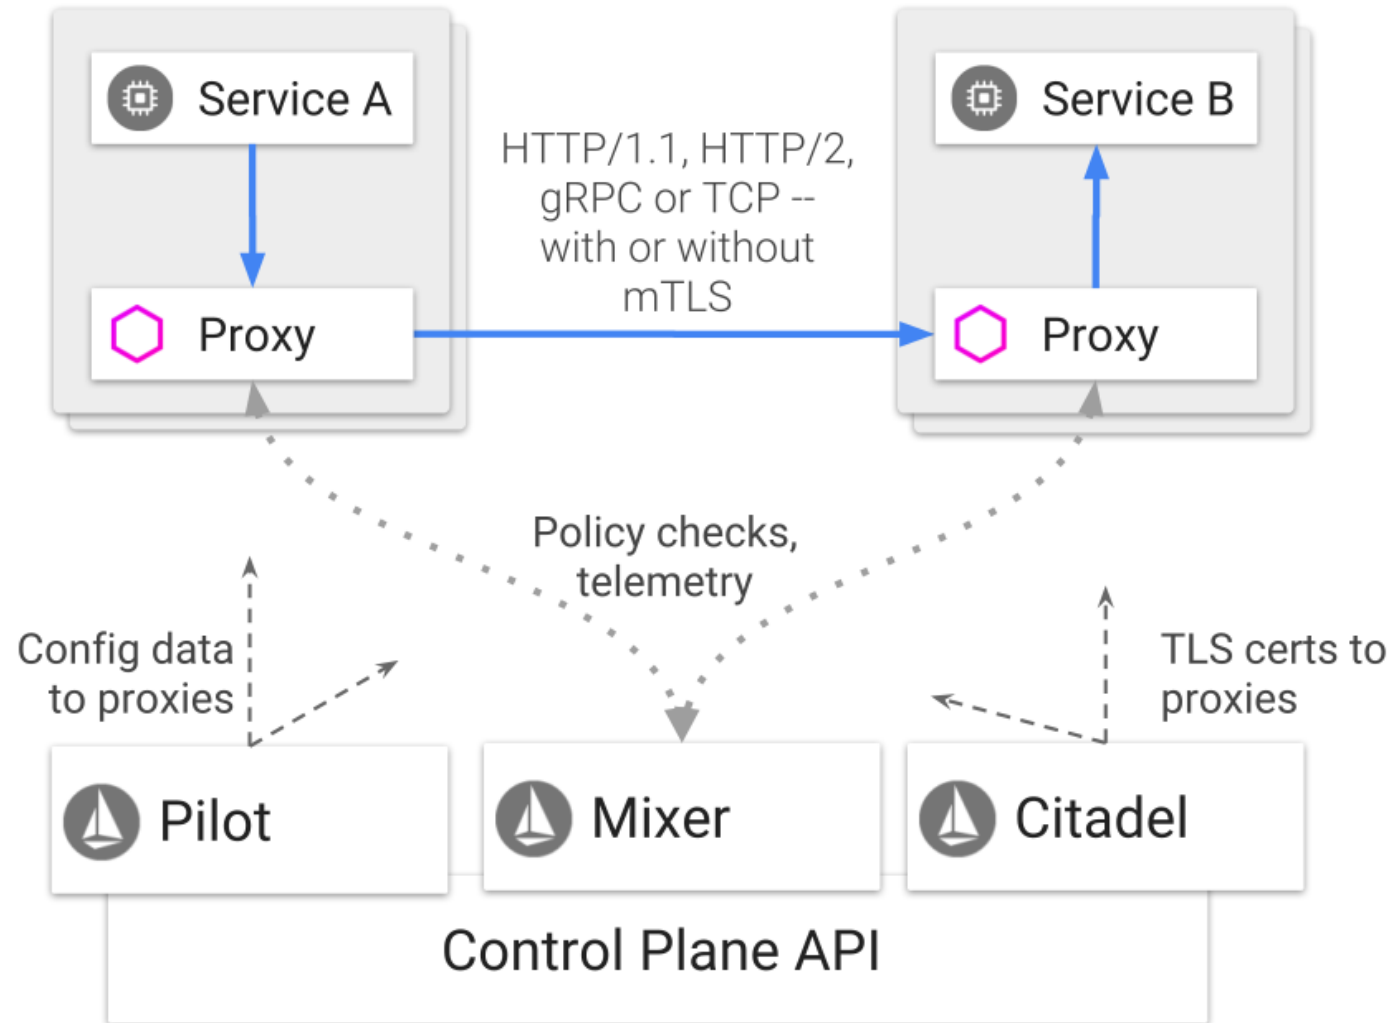


# What is Istio?

- Istio is a service mesh platform providing
  - Observability
  - Traffic management
  - Security
- Install on Kubernetes as a set of Pod and services to `istio-system` namespace



# Istio Architecture





# Istio Architecture

## Data Plane

- A set of sidecars (proxies)
- Sidecars functionalities
  - Dynamic service discovery
  - Client side load balancing
  - TLS termination
  - HTTP/2 and gRPC proxies
  - Circuit breakers
  - Health checks
  - Metrics
  - Weighted traffic split

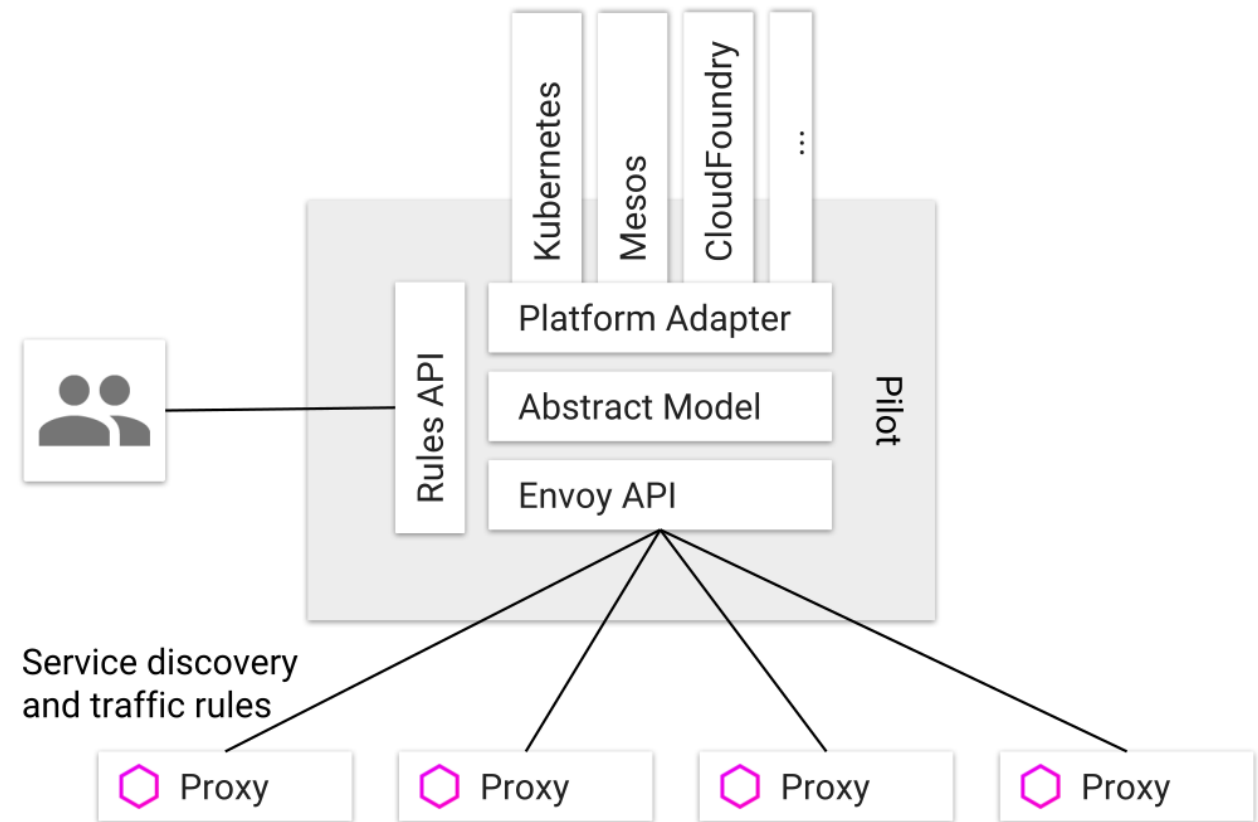
## Control Plane

- Manages and configures the sidecars
- Enforces policies
- Has the following components
  - Mixer enforces access control and usage policies across the service mesh
  - Pilot provides service discovery for the sidecars
  - Citadel provides service to service and end-user authentication



# Pilot and Proxy

- Pilot is part of Istio's control plane
- Use to configure the proxy in each of the Pod
  - Timeout information
  - How to handle failures
  - Routing policies
- Proxy enforces the configurations and policies from Pilot







# Installing Istio - with Helm

- Download and unzip/untar Istio
  - <https://github.com/istio/istio/releases>
- Create custom resources
  - In `install/kubernetes/helm` directory

```
helm install istio-init \  
  --name istio-init \  
  --namespace istio-system
```
- Create Istio resources
  - See <https://istio.io/docs/setup/kubernetes/additional-setup/config-profiles> for list of installed components

```
helm install istio \  
  --values istio/values-istio-demo.yaml
```



# Istio Pods and Services

- Istio installs its artefacts into `istio-system` namespace
- Pods include (not exhaustive)
  - `istio-citadel`
  - `istio-ingressgateway`
  - `istio-egressgateway`
  - `istio-tracing`
- Services include (not exhaustive)
  - `istio-egressgateway`
  - `istio-ingressgateway` - traffic enters the cluster via this
  - `jaeger-query`



# Deploying to Istio

- Create a deployment and service object as per normal
- Except inject sidecar into Pods

```
istioctl kube-inject -f deployment.yml > \
    deployment-with-sidecar.yml
kubectl apply -f deployment-with-sidecar.yml
```

or

```
kubectl apply -f \
    <(istioctl kube-inject -f deployment.yml)
```

- Deploy service

```
kubectl apply -f service.yml
```

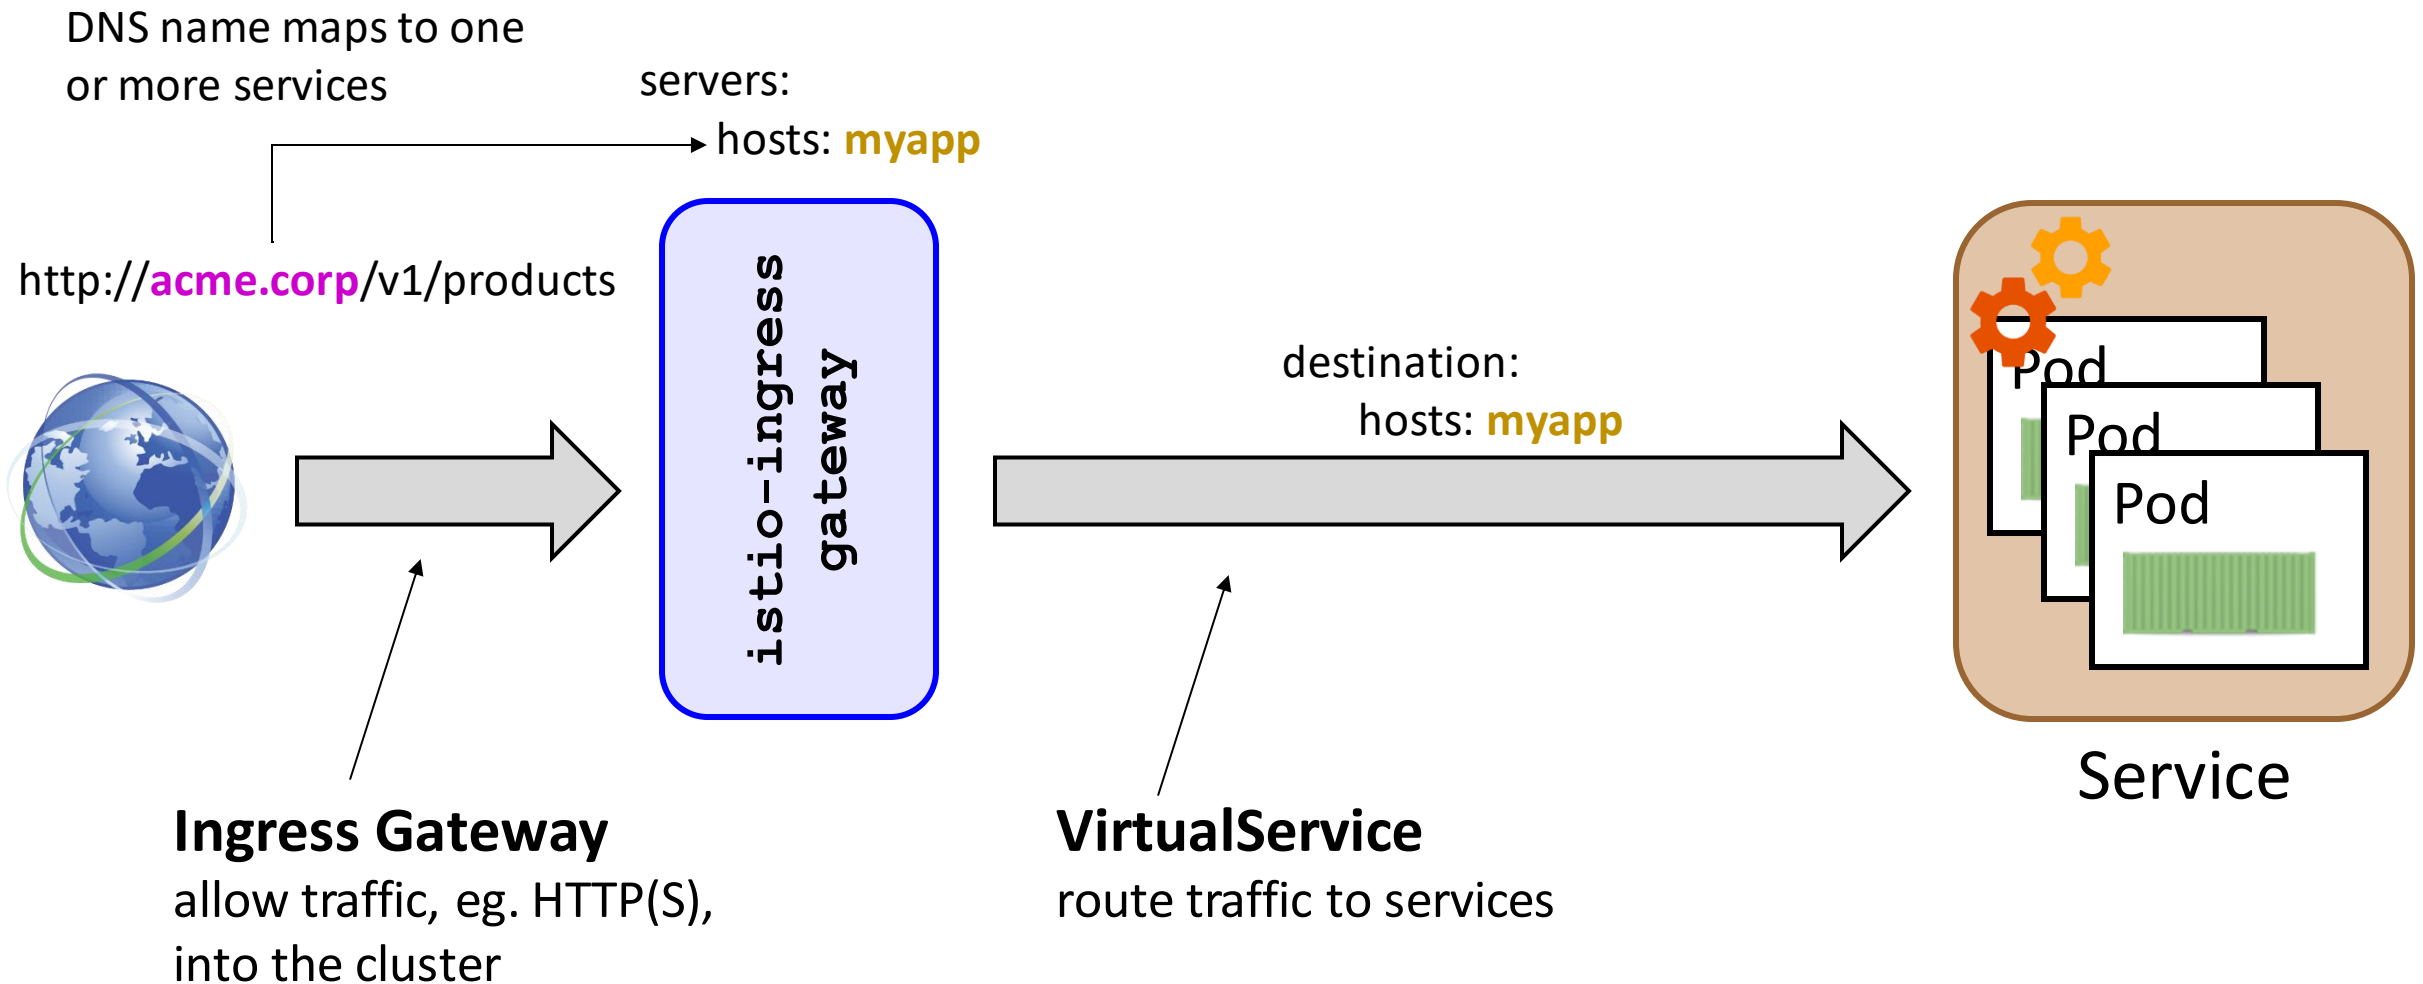


# Istio Routing

- Uses Gateway and VirtualService inplace of Ingress
  - Gateway L4 to L6 load balancer
    - Configures sidecar for ingress
  - VirtualService is a L7 router
    - Bound to a gateway, route traffic to services
- Decoupling of traffic with Gateway and VirtualService provides different traffic management features configured at the VirtualService
  - Eg. A/B testing, canary releases, gradual rollouts, etc.
- Routing can be applied to both ingress and egress traffic



# Gateway and VirtualService - Ingress





# Define a Gateway

```
apiVersion: networking.istio.io/v1alpha3
```

```
kind: Gateway
```

```
metadata:
```

```
  name: myapp-gateway
```

```
spec:
```

```
  selector:
```

```
    istio: ingressgateway
```

```
  servers:
```

```
    - port:
```

```
      number: 80
```


```
      name: http
```

```
      protocol: HTTP
```

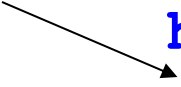
```
    hosts:
```

```
      - myapp.ns.svc.cluster.local
```

Use Istio ingress to  
handle this



Traffic to this host  
(service) or \* for  
any service. Must  
be FQDN



Accept traffic to this gateway  
from one or more ports





# Define a Virtual Service

```
apiVersion: networking.stio.io/v1alpha3
```

```
kind: VirtualService
```

```
metadata:
```

```
  name: myapp
```

```
spec:
```

```
  hosts:
```

```
  - "*" ← Receive traffic from the
```

```
  gateways:
```

```
  - myapp-gateway ← list of gateways
```

```
  http:
```

```
  - match:
```

```
    - uri:
```

```
      exact: /v1 ← All request with
```

literal match

```
    route:
```

```
    - destination:
```

```
      hosts: myapp.ns.svc.cluster.local ← Route to this service.
```

Must be FQDN

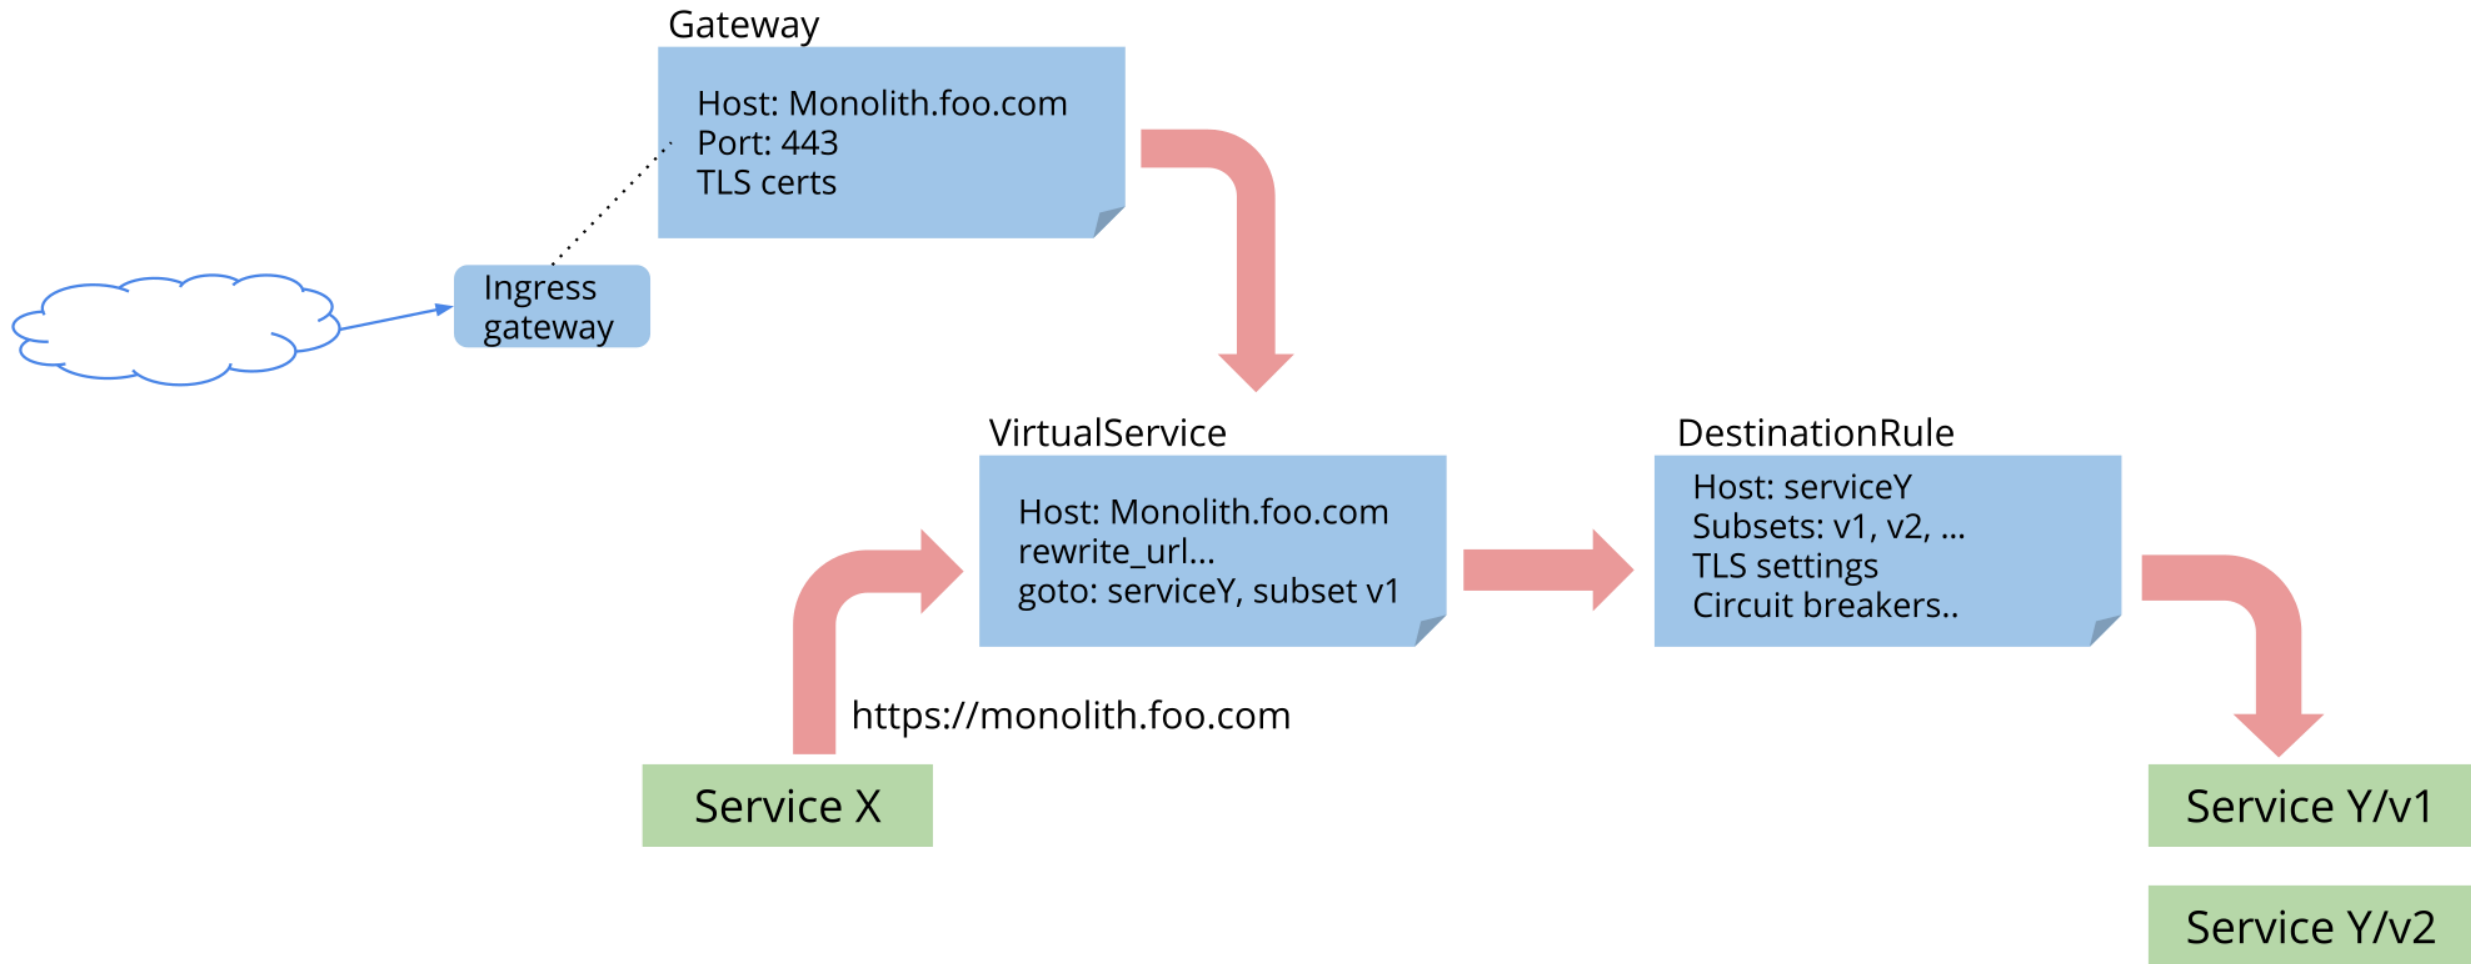
Rewrite the resource  
from /v1 to /

```
    rewrite:
```

```
      uri: /
```



# Istio Routing







# Default Route

```
apiVersion: networking.stio.io/v1alpha3
kind: VirtualService
metadata:
  name: myapp
spec:
  hosts:
  - "*"
  gateways:
  - myapp-gateway
  http:
  ...
```

Request with /v1 are routed to myapp, all other request defaults to myapp\_v2

One or more matched route

Default route

```
  http:
  - match:
    - uri:
        exact: /v1
    route:
    - destination:
        hosts: myapp
      rewrite:
        uri: /
  - route:
    - destination:
        hosts: myapp_v2
        port:
          number: 8080
```



# Weight Based Routing

```
apiVersion: networking.stio.io/v1alpha3
kind: VirtualService
spec:
  gateways:
  - myapp-gateway
  http:
  - match:
    - uri:
        exact: /v1
    route:
    - destination:
      hosts: myapp
      weight: 75
    - destination:
      hosts: myapp-beta
      weight: 25
```

If there is only 1 destination, the weight property can be omitted

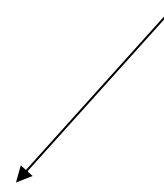
Traffic is split between the 2 services. Assume myapp-beta is a new version that we are testing



# Routing Based on HTTP Header

```
apiVersion: networking.stio.io/v1alpha3
kind: VirtualService
spec:
  hosts:
  - myapp
  gateways:
  - myapp-gateway
  http:
  - match:
    - headers:
      authorization: ^.*Bearer .+$
    route:
      - destination:
          hosts: myapp
```

Route bases on the value  
of HTTP headers





# Routing Based on Pod Labels

```
apiVersion: networking.stio.io/v1alpha3
kind: VirtualService
spec:
  hosts:
  - myapp
  gateways:
  - myapp-gateway
  http:
  - match:
    - sourceLabels:
      app: myotherapp
      version: v2
    route:
    - destination:
        hosts: myapp
```

Only allow request from Pods  
with the following labels





# Configure CORS

```
apiVersion: networking.stio.io/v1alpha3
```

```
kind: VirtualService
```

```
spec:
```

```
  hosts:
```

```
  - myapp
```

```
  gateways:
```

```
  - myapp-gateway
```

```
  http:
```

```
  - match:
```

```
    - headers:
```

```
      authorization: ^.*Bearer .+$
```

```
    route:
```

```
    - destination:
```

```
      hosts: myapp
```

```
  corsPolicy:
```

```
    allowOrigin:
```

```
    - *
```

```
    allowMethods:
```

```
    - GET
```

```
    - POST
```

Returns the following CORS headers in the response

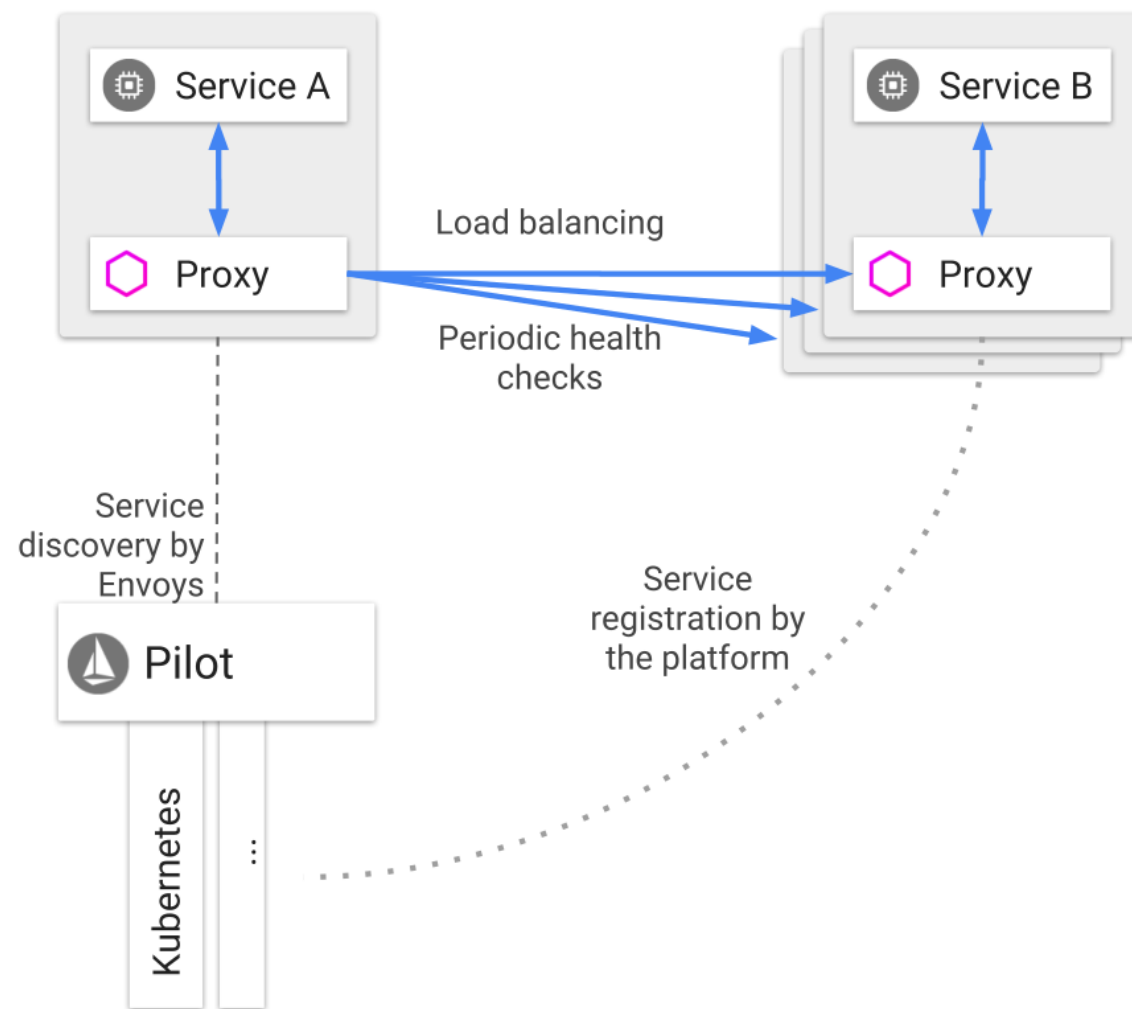
Access-Control-Allow-Origin: \*

Access-Control-Allow-Methods: GET, POST



# Load Balancing and Health Checks

- Istio Pilot gets service information from Kubernetes
- Proxy (sidecar) discovers these services from Pilot
- Proxy keeps a list of the pods in the services internally
- Proxy load balances across the Pods in a service
- Proxy also does health checks on the Pods
  - Removes unhealthy Pods from its lists





# Define a DestinationRule

- Applies to traffic bound for a particular service
  - After routing has occurred

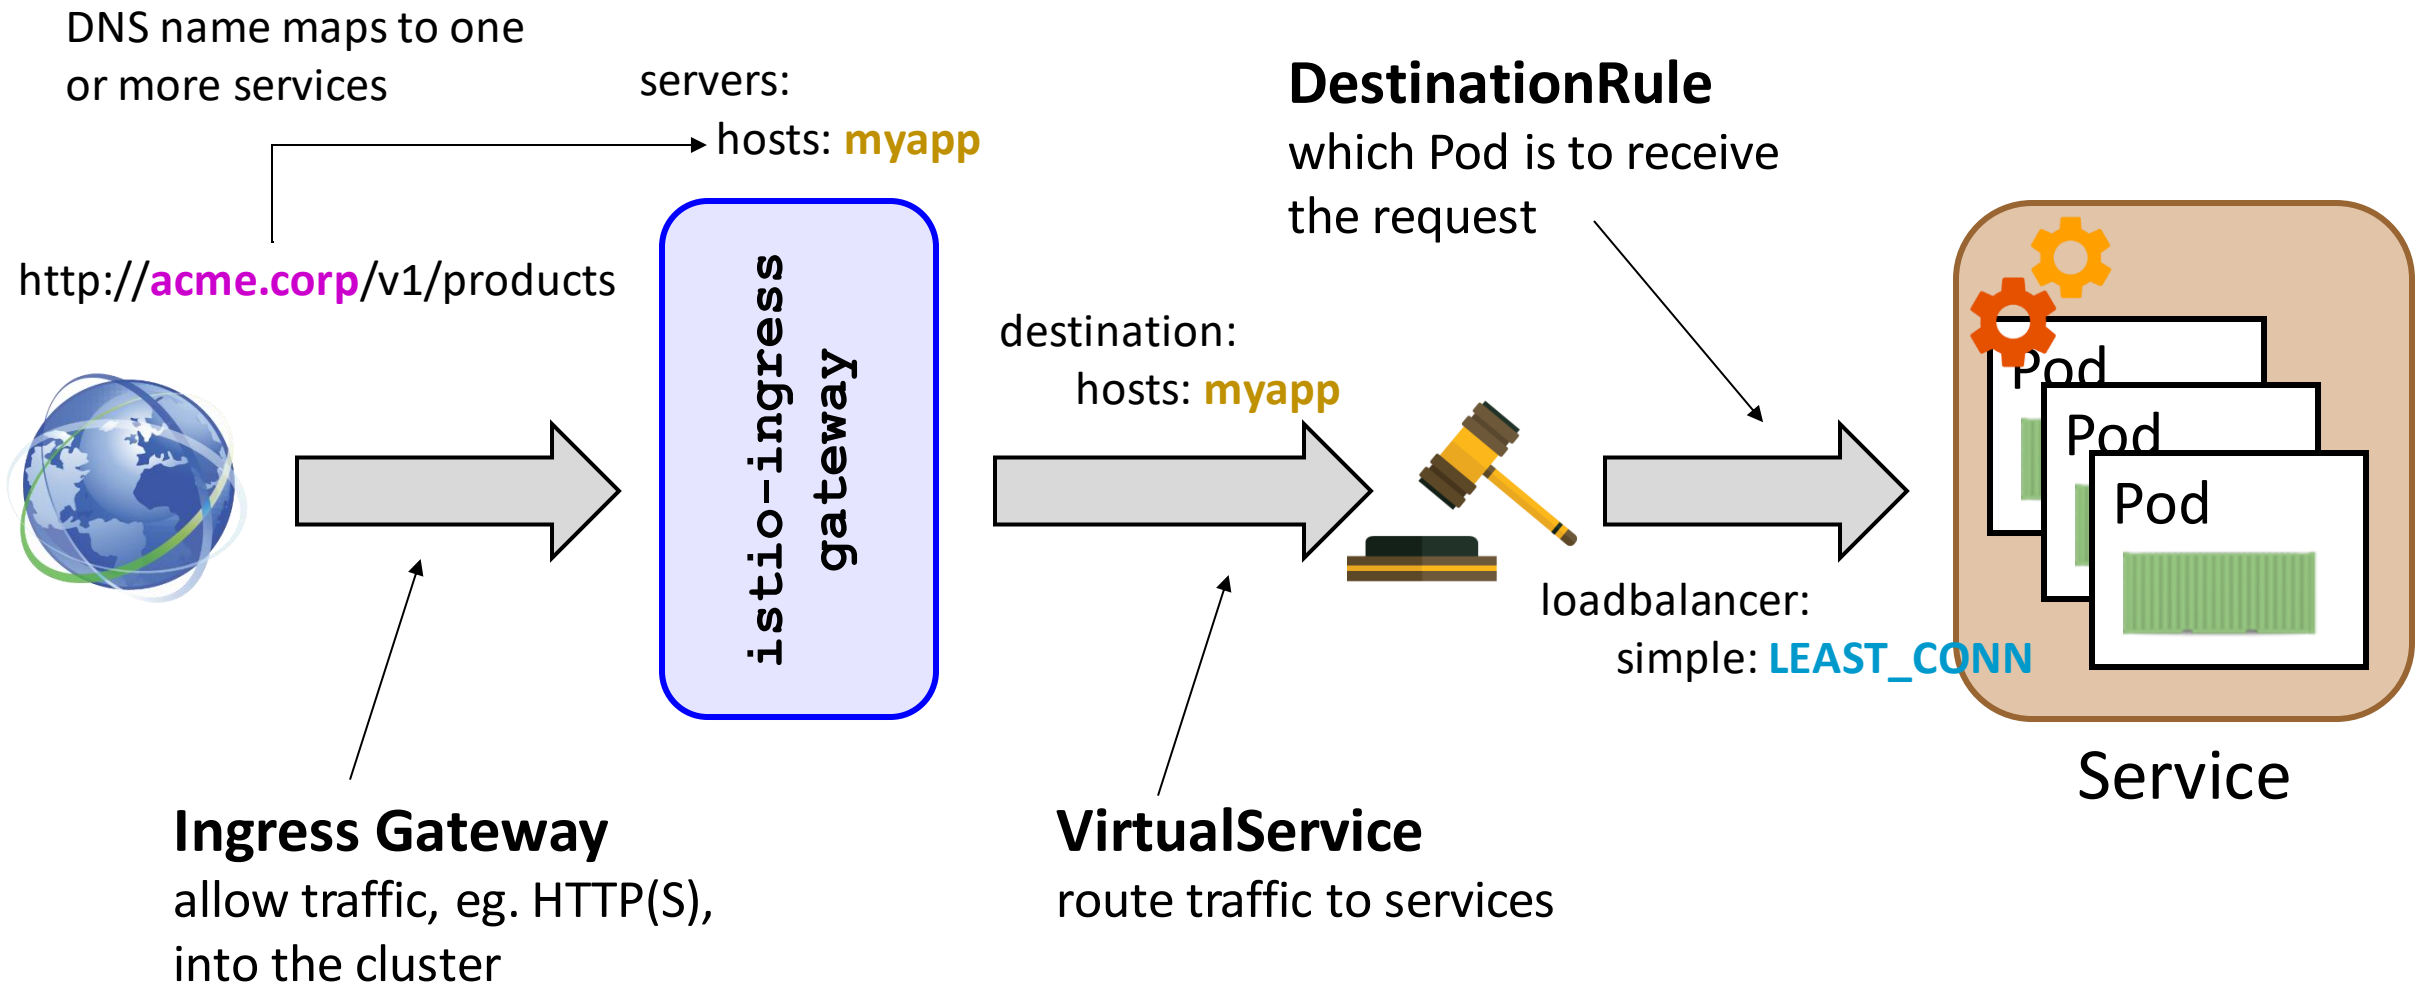
```
apiVersion: networking.istio.io/v1alpha3
kind: DestinationRule
metadata:
  name: myapp
spec:
  host: myapp
  trafficPolicy:
    loadBalancer:
      simple: LEAST_CONN
```

## Load balancer policy

- ROUND\_ROBIN
- LEAST\_CONN - least busy host
- RANDOM - randomly selects a healthy host
- PASSTHROUGH - pass directly to the requested host



# Gateway and VirtualService - Ingress








# VirtualService with DestinationRule - 2

```
apiVersion: networking.stio.io/v1alpha3
kind: VirtualService
metadata:
  name: myapp
spec:
  hosts:
  - myapp
  gateways:
  - myapp-gateway
  http:
  - match:
    - uri:
        exact: /v1
    route:
    - destination:
        hosts: myapp
        subset: v1rule
```

Destination  
rule name

An arrow pointing from the text "Destination rule name" to the value "v1rule" in the subset field of the destination rule.



# VirtualService with DestinationRule - 2

```
apiVersion: networking.istio.io/v1alpha3
kind: DestinationRule
metadata:
  name: myapp
spec:
```

```
  host: *
```

```
  subsets:
```

```
    - name: v1rule
```

```
      labels:
```

```
        app: myapp
```

```
        version: v1
```

```
      trafficPolicy:
```

```
        loadBalancer:
```

```
          simple: LEAST_CONN
```

Destination rule name

Apply to Pods with these labels

One or more  
named  
destination  
rule



# VirtualService Route Matching Example

Ingress Resource	Rewrite To
<code>/v1</code>	<code>/</code>
<code>/v1/</code>	<code>/</code>
<code>/v1/products</code>	<code>/products</code>
<code>/v1/products/tv</code>	<code>/products/tv</code>

Order of the match is important

```
http:
- match:
  - uri:
      exact: /v1
    rewrite:
      uri: /
    route:
      - destination:
          host: myapp
  - match:
    - uri:
        prefix: /v1
      rewrite:
        uri: " "
      route:
        - destination:
            host: myapp
```



# Egress

`fetch()` made from within a container in a Pod

Sidecar will attempt to resolve this name.  
Not found in Pilot's internal registry

```
fetch('https://hacker-news.firebaseio.com/v0/item/1234.json')  
  .then(response => {
```

- All traffic from container(s) within a Pod are routed through the sidecar
- The sidecar routes the traffic to other services via service discovery
- Pod will not be able to make external calls viz. to resources outside of the service mesh



# Handling Egress Traffic

- Two ways to allow traffic out of Istio
  - Creating a service entry to allow individual Pods to access to destination
  - Going through an egress gateway



# Service Entry

- Manually add external resources into the Istio's service registry
  - Maintained by Pilot
- Allow sidecar to discover the service
- Describes the service characteristics
  - Eg. host, ports, TLS, etc

**apiVersion:** networking.istio.io/v1alpha3

**kind:** ServiceEntry

metadata:

name: myapp-egress

spec:

**hosts:**

- hacker-news.firebaseio.com
- api.openweathermap.org

List of external host resolvable by DNS

**ports:**

- **number:** 80  
**protocol:** HTTP  
**name:** http
- **number:** 443  
**protocol:** HTTPS  
**name:** https

One or more ports associated with the external service

**location:** MESH\_EXTERNAL

The service is outside of the mesh

**resolution:** DNS

Mode in which the service is resolved



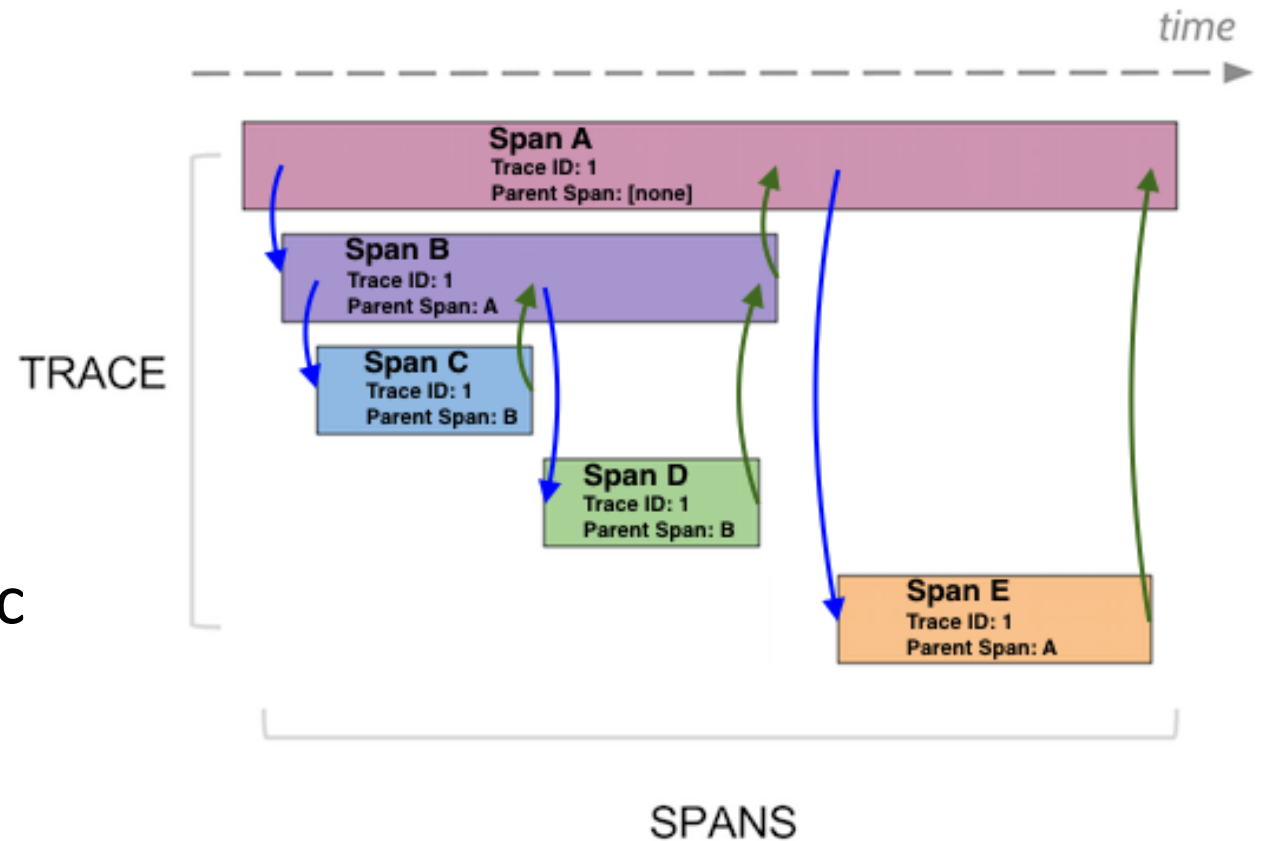
# Distributed Tracing

- To follow the request as it traverses the application mesh
  - Provides observability as a request is being fulfilled by various micro services
- Captures the entire call chain
  - From when the request enters the service mesh and all the services that the request interacts with
- For monitoring and profiling micro services
- OpenTracing is an open standard for distributed tracing
  - Jaeger is an implementation of OpenTracing



# OpenTracing

- Trace covers the entire request across all services that the request pass thru
- Trace consist of one or more Spans
- Span is a logical chunk of work
  - Spans can contain other spans
- Micro services propagate specific HTTP headers to outgoing request to enable tracing
  - Eg. X-B3-TraceId, X-B3-SpanId, X-B3-Sampled, etc.







# Tracing Console

- Find the distributed tracing pod, starts with `istio-tracing-XXXX`

```
kubectl get pod -n istio-system | grep tracing
```

- Port forward the console

```
kubectl port-forward istio-tracing-XXX 16686:16686
```



# Example of a Trace





# Appendix



# Installing Istio - with kubectl

- Download and unzip/untar Istio
  - <https://github.com/istio/istio/releases>

- Create custom resources

```
kubectl apply -f  
<istio_home>/install/kubernetes/helm/istio/templates/crds.  
yaml
```

- Create Istio resources

- Without authentication

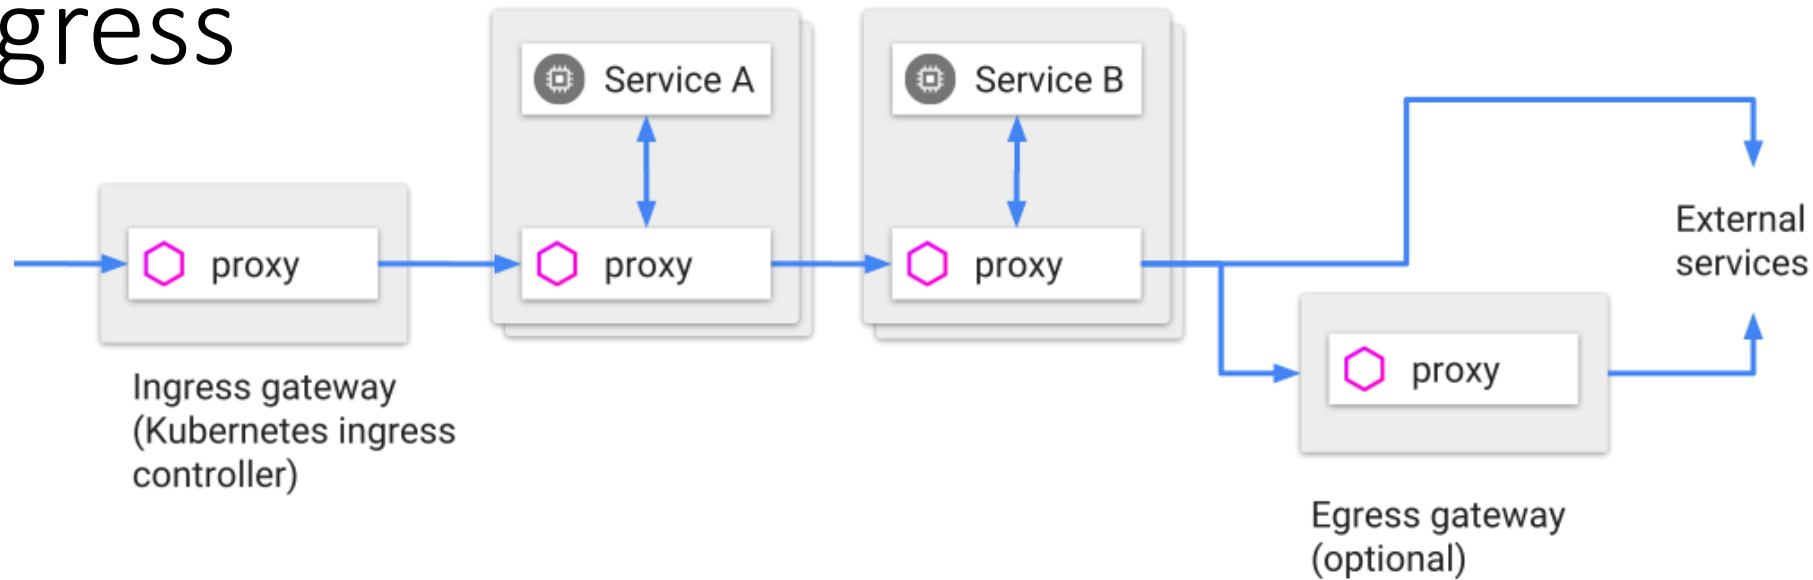
```
kubectl apply -f <istio_home>/install/kubernetes/istio-  
demo.yaml
```

- With authentication

```
kubectl apply -f <istio_home>/install/kubernetes/istio-  
demo-auth.yaml
```



# Egress



- Egress connections for sidecar is forwarded to an egress gateway
- Egress gateway allows for greater control eg.
  - Logging
  - Converting all non TLS traffic to TLS
  - Cache results
  - Circuit breaking



# Routing to External Service - Gateway

```
apiVersion: .../v1alpha3
kind: Gateway
metadata:
  name: myapp-egress
spec:
  selector:
    istio: egressgateway
  servers:
  - hosts:
    - api.openweathermap.org
    - hacker-news.firebaseio.com
    port:
      number: 443
      protocol: HTTPS
      name: https
    tls:
      mode: SIMPLE
      serverCertificate: /etc/certs/cert-chain.pem
      privateKey: /etc/certs/key.pem
      caCertificates: /etc/certs/root-cert.pem
```

Default list of certificate in the egressgateway Pod. To examine the certs, find the egressgateway Pod name in istio-system namespace and exec into it

} SSL connections for hosts



# Routing to External Service - VirtualService

```
apiVersion: .../v1alpha3
kind: VirtualService
metadata:
  name: myapp-egress
spec:
```

```
  hosts:
```

```
    - api.openweathermap.org
```

```
  gateways:
```

```
    - mesh
```

```
    - myapp-egress
```



List of gateways  
to use

```
  http:
```

```
    - match:
```

```
      - gateways:
```

```
        - mesh
```

```
        port: 80
```

```
      route:
```

```
        - destination:
```

```
          host: istio-egress-istio-system.svc.cluster.local
```

Traffic going to api.openweathermap.org is  
required to use myapp-egress gateway

```
    - match:
```

```
      - gateways:
```

```
        - myapp-egress
```

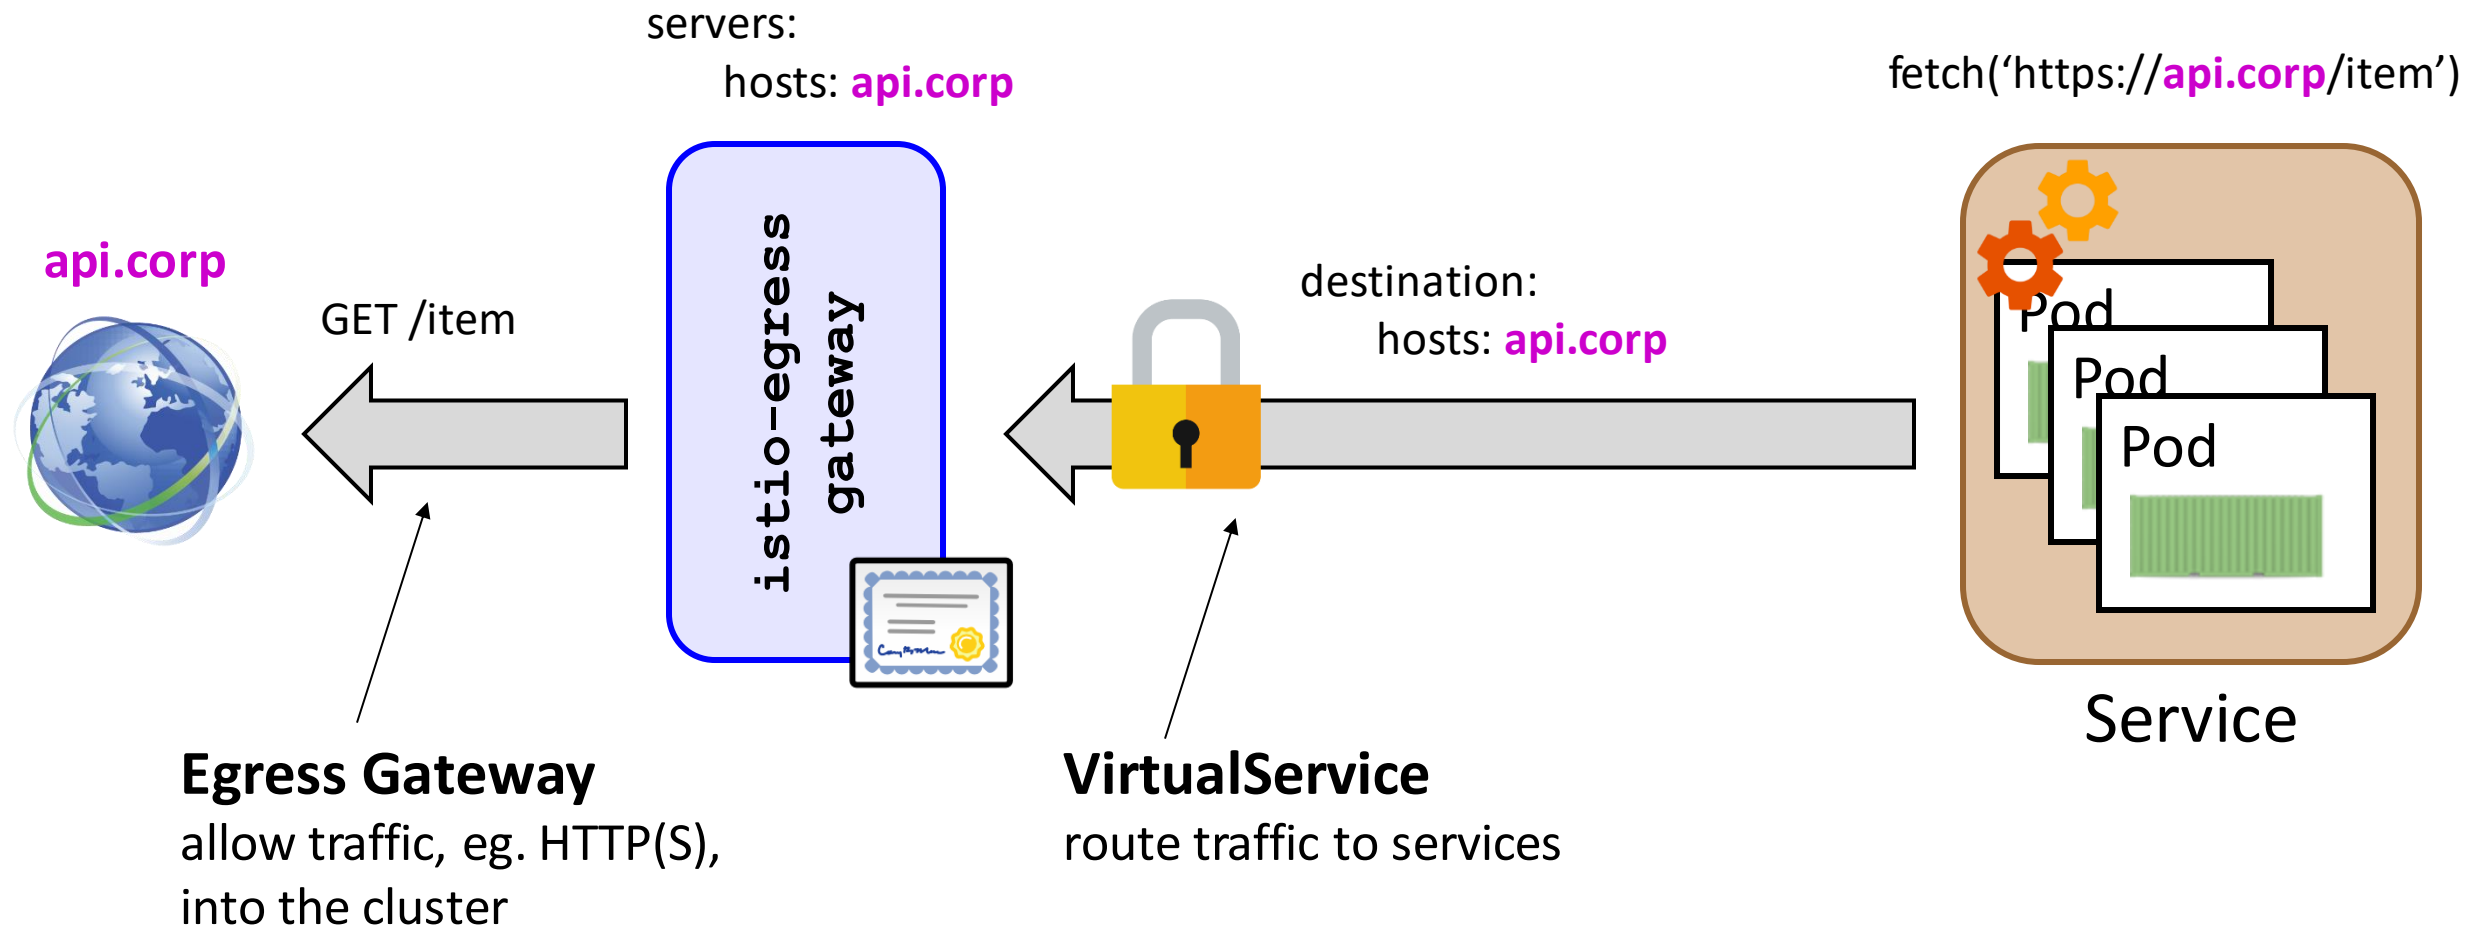
```
        port: 80
```

```
      route:
```

```
        - destination:
```

```
          host: api.openweathermap.org
```

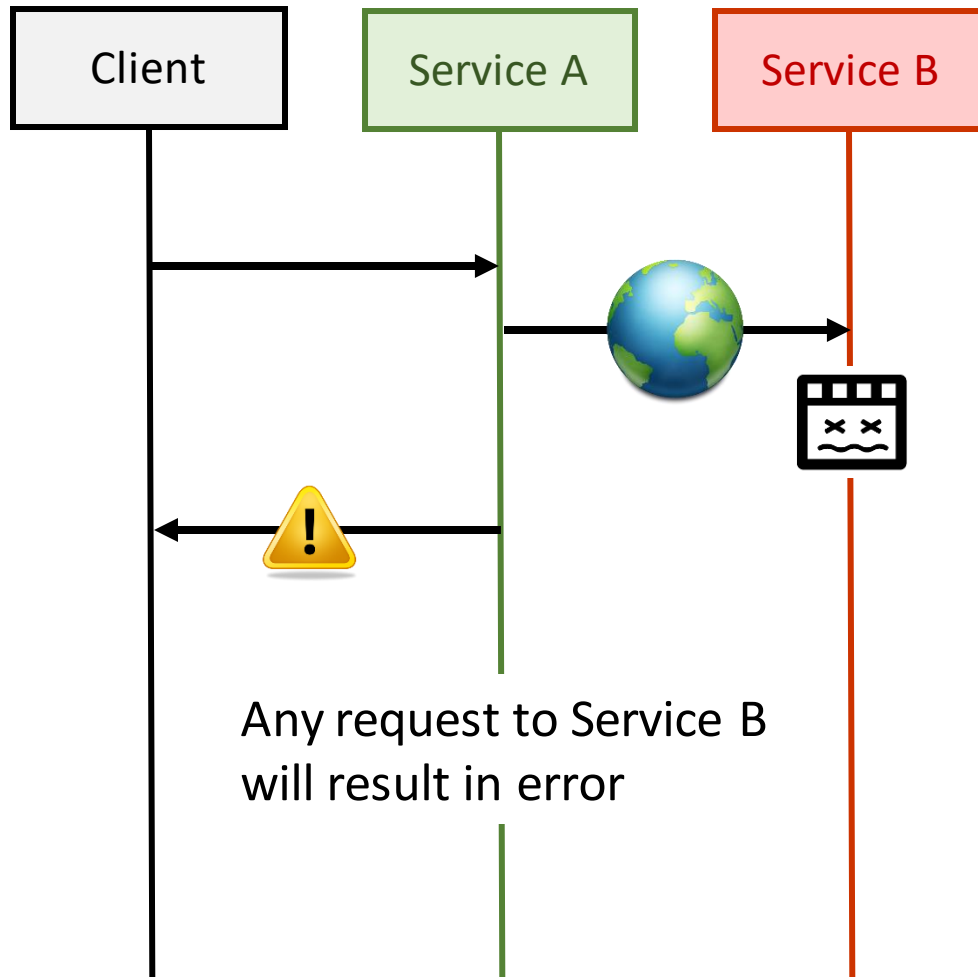
# Gateway and VirtualService - Egress







# Circuit Breaker

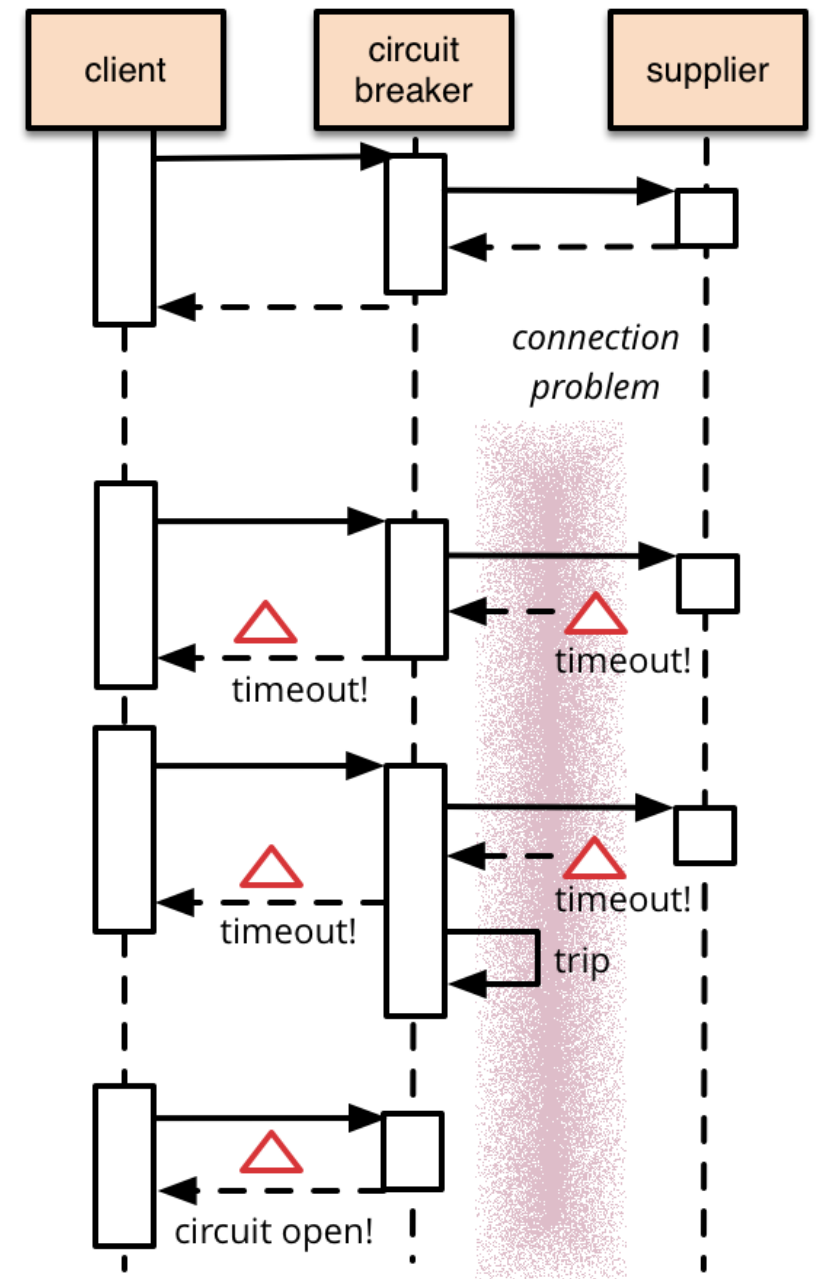


- Lots of things can go wrong when connecting to a remote connection
  - Not under your control
  - Eg. network issue, service down, etc.
- If Service B is down, then there is no point in making any further calls to it until it becomes available
  - Continually make calls from Service A may lead to cascading failures



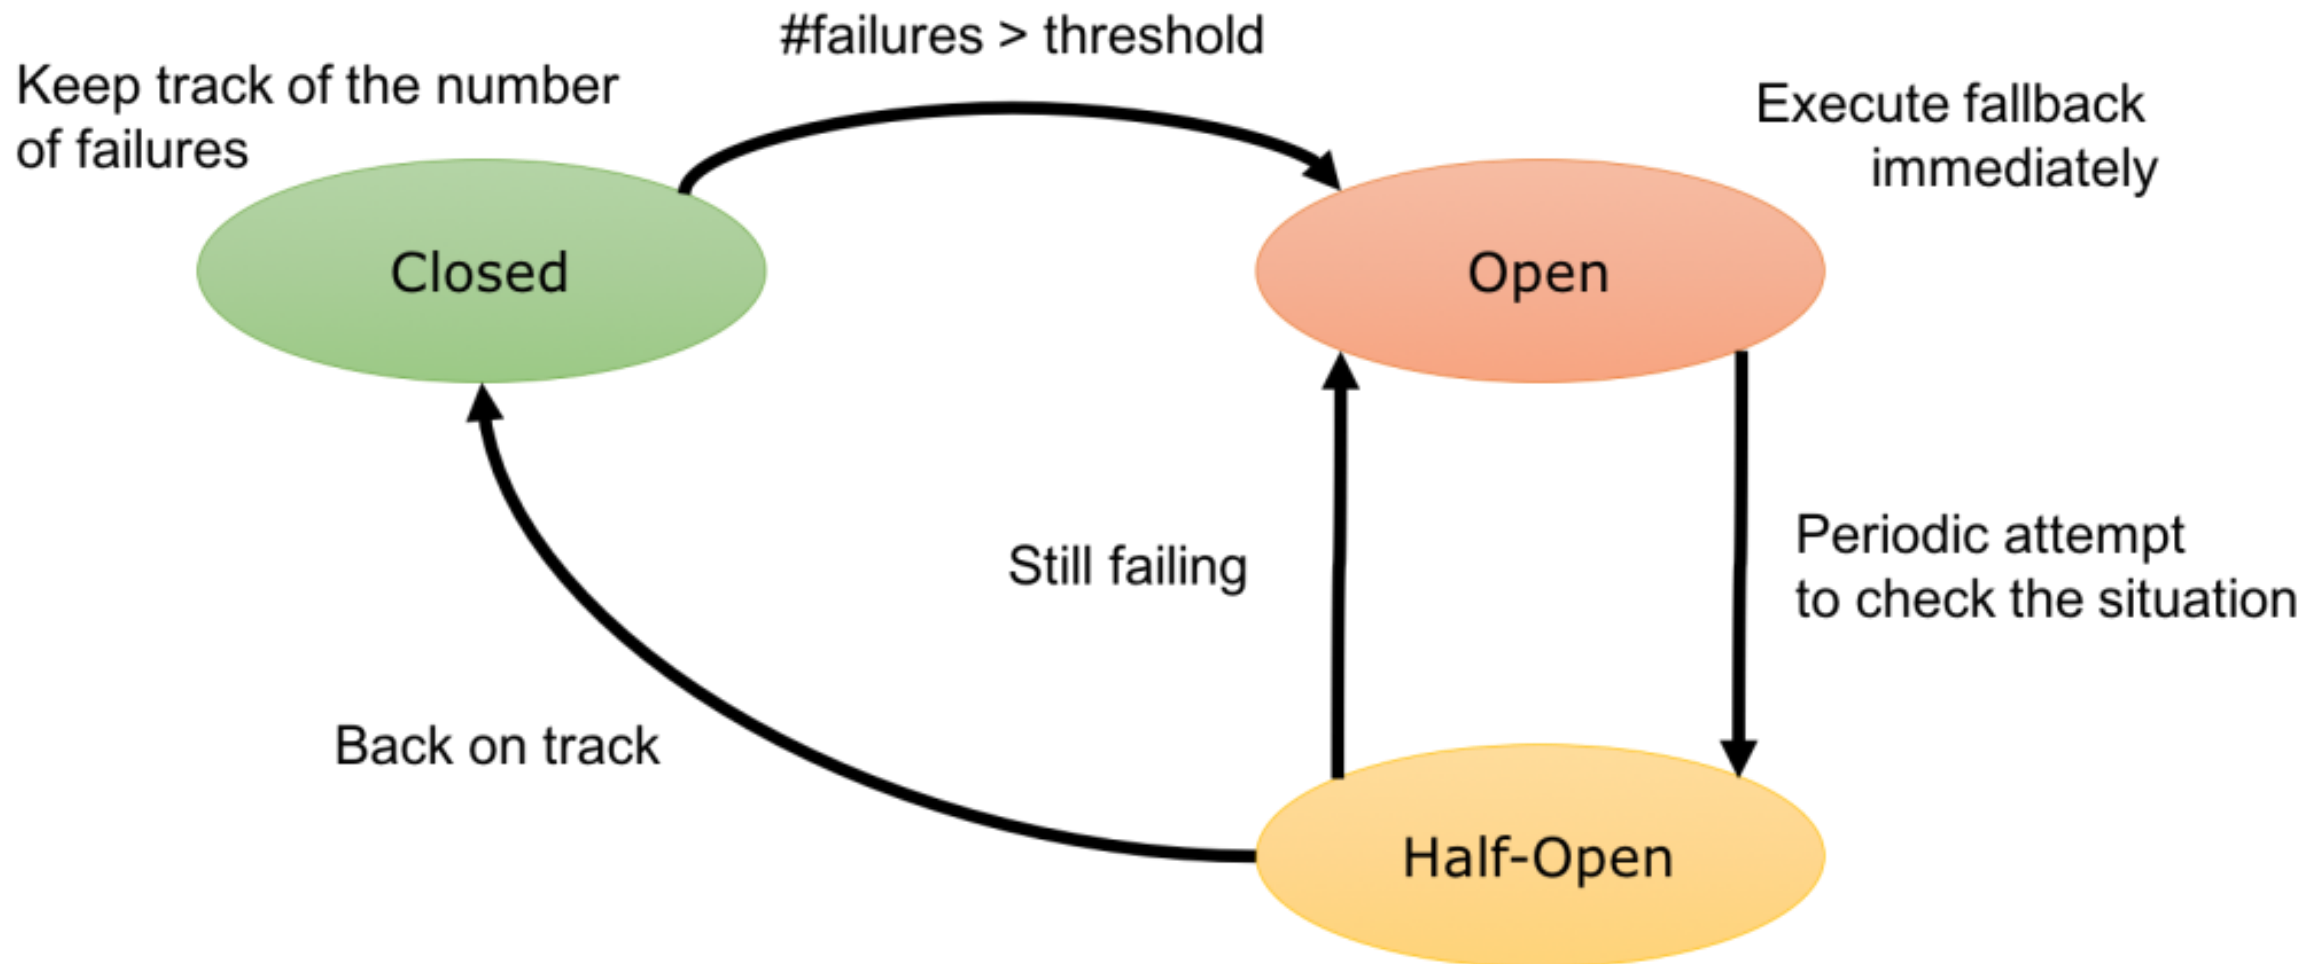
# Circuit Breaker

- Remote calls are proxied by a circuit breaker object
- Circuit breaker monitors the availability of remote services
- Once a failure threshold is breached, the circuit is open
- Any subsequent call to the service via the circuit breaker will result in an immediate error





# Circuit Breaker Algorithm





# Defining a Circuit Breaker

```
apiVersion: networking.istio.io/v1alpha3
kind: DestinationRule
metadata:
  name: myapp-egress
spec:
  host: api.openweathermap.org
  trafficPolicy:
    connectionPool:
      maxConnection: 10
    outlierDetection:
      interval: 5m
      consecutiveErrors: 5
      baseEjectionTime: 15m
```

Destination

Optional. Number of  
egress connection in  
the gateway

Number of consecutive  
5XX errors in a period of 5  
minutes, the destination  
will be ejected for 15  
minutes